

High performance parallel systems


Lecture 1 – Basics of Computer Architecture

Kenneth Skovhede, NBI, 2020-11-17

UNIVERSITY OF COPENHAGEN



Analog signals




Analog signals

- have (theoretically) unlimited resolution
- invariably degrade (noise)
- accumulate noise (signal-to-noise ratio increases)
- eventually lose the signal in noise



Digital signals

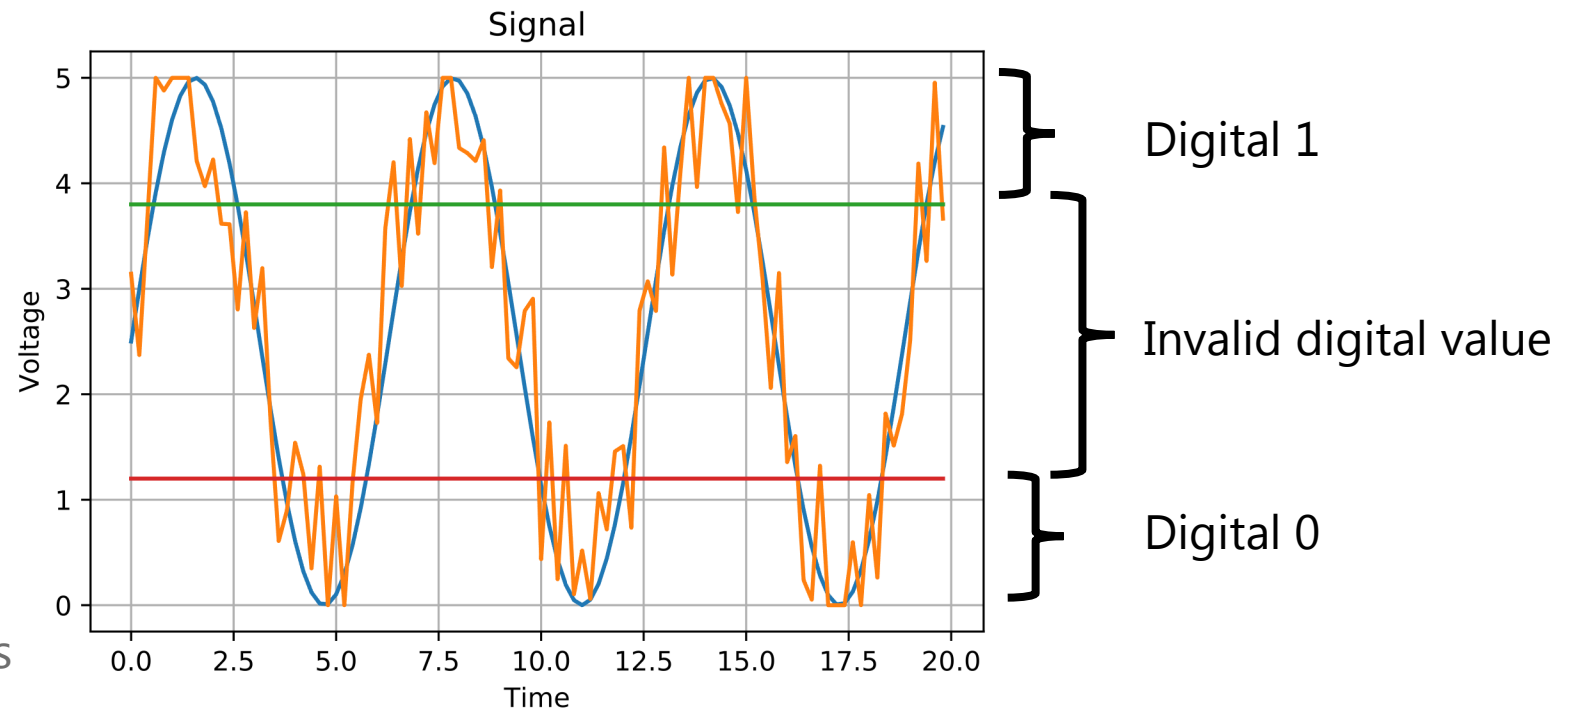



Digital signals

- have on/off resolution
- tolerate degradation (noise)
- removes noise when forwarded
- can be repeated infinitely (theoretically)
 - (bit flips occur, but for other reasons)



Digital signals



We can use other interpretations than binary but it generally just increases complexity



A game of telephone

Unlimited vocabulary



Binary vocabulary





Numbers with a digital system



Digital and binary

Binary (0b)	Decimal	Binary (0b)	Decimal
0000	0	1011	11
0001	1	1100	12
0010	2	1101	13
0011	3	1110	14
0100	4	1111	15
0101	5	10000	16
0110	6	10001	17
0111	7	10010	18
1000	8	10011	19
1001	9	10100	20
1010	10





Try it yourself – base2 & base4 addition

To get a feel for math with numbers that do not use base10, try these basic math problems:

Binary numbers:	Result (in binary)
<u>0b0010</u> + 0b0010 =	
0b0101 + 0b1010 =	
0b0111 + 0b0111 =	

Base4 numbers:	Result (in base4)
0010 + <u>0010</u> =	
0301 + 1030 =	
0333 + 0333 =	

Do not convert into decimal before doing the addition!





Try it yourself – base2 & base4 addition

To get a feel for math with numbers that do not use base10, try these basic math problems:

Binary numbers:	Result (in binary)
$0b0010 + 0b0010 =$	$0b0100$
$0b0101 + 0b1010 =$	$0b1111$
$0b0111 + 0b0111 =$	$0b1110$



Base4 numbers:	Result (in base4)
$0010 + 0010 =$	0020
$0301 + 1030 =$	1010
$0333 + 0333 =$	1332

Do not convert into decimal before doing the addition!



Other numeric bases


Binary (0b)	Decimal	Octal (0c)	Hexadecimal (0x)
0000	0	0	0
0001	1	1	1
0010	2	2	2
0011	3	3	3
0100	4	4	4
0101	5	5	5
0110	6	6	6
0111	7	7	7
1000	8	10	8
1001	9	11	9
1010	10	12	A
1011	11	13	B
1100	12	14	C
1101	13	15	D
1110	14	16	E
1111	15	17	F
10000	16	20	10



Try it yourself – base16 addition

To get a feel for math with numbers that do not use base10, try these basic math problems:

Hex numbers:	Result (in hex)
$\text{0xA} + \text{0x1} =$	
$\text{0x1F} + \text{0x1F} =$	
$\text{0xBCBC} + \text{0x1234} =$	




Do not convert into decimal before doing the addition!



Try it yourself – base16 addition

To get a feel for math with numbers that do not use base10, try these basic math problems:

Hex numbers:	Result (in hex)
$0xA + 0x1 =$	$0xB$
$0x1F + 0x1F =$	$0x3E$
$0xBCBC + 0x1234 =$	$0xCEF0$



Do not convert into decimal before doing the addition!



Bases with "power of 2" properties

Binary	Decimal	Octal	Hexadecimal
0000	0	0	0
0001	1	1	1
0010	2	2	2
0011	3	3	3
0100	4	4	4
0101	5	5	5
0110	6	6	6
0111	7	7	7
1000	8	10	8
1001	9	11	9
<u>1010</u>	10	12	<u>A</u>
<u>1011</u>	11	13	<u>B</u>
1100	12	14	C
1101	13	15	D
1110	14	16	E
1111	15	17	F

Hexadecimal	Binary
<u>0xAB</u>	<u>1010 1011</u>
0xF1	1111 0001
<u>0xABCDEF09</u>	<u>1010 1011 1100 1101 1110 1111 0000 1001</u>

Octal	Binary
0c56	101 110
0c71	111 001
0x6701	110 111 000 001




Negative numbers

Bits	Min	Max
4	-8	7
8	-128	127
16	-16384	16383
32	-2147483648	2147483647

Odd behavior due to non-symmetry


- $-(-8) = -8$
- $\text{Abs}(-8) = -8$
- $-8 * -1 = -8$



Digital circuits



Building blocks – A transistor



Note that the image on the right is the transistor with a plastic house, which is obviously not present inside a CPU core



Basic gates

AND	0	1
0	0	0
1	0	1

AND



OR




OR	0	1
0	0	1
1	1	1

XOR	0	1
0	0	1
1	1	0



XOR



Not

NOT	0	1
	1	0



A conceptual computer



The Turing Machine





Image from: <https://phys.org/news/2013-03-artificial-muscle-universal-turing-machine.html>



Turing Machine Example

- The state can be any of: A, B, C, ~~D~~
 - The state D means program termination
 - Initialized to state A
- The tape can contain either zero or one
 - The tape is initialized to all zeros
- Write a program that emits the sequence 101010...

Input	Action
A0	<u>B1 R</u>
<u>B0</u>	<u>A0 R</u>
C0	unreachable
A1	unreachable
B1	unreachable
C1	unreachable

New state = B, Write 1, Move right

New state = A, Write 0, Move right

Tape state

0	1	2	3	4
0	0	0	0	0


Head



Try it yourself – Turing Machine

- The state can be any of: A, B, C, D
 - The state D means program termination
 - Initialized to state A
- The tape can contain either zero or one
 - The tape is initialized to all zeros
- Write a program that emits the sequence 1100000...


Input	Action
A0	
B0	
C0	
A1	
B1	
C1	



Try it yourself – Turing Machine

- The state can be any of: A, B, C, D
 - The state D means program termination
 - Initialized to state A
- The tape can contain either zero or one
 - The tape is initialized to all zeros
- Write a program that emits the sequence 1100000...

Input	Action
A0	B1, R
B0	C1, R
C0	C0, R
A1	
B1	
C1	



The halting problem

```
void easy_halting1(int x) {  
    return;  
}
```

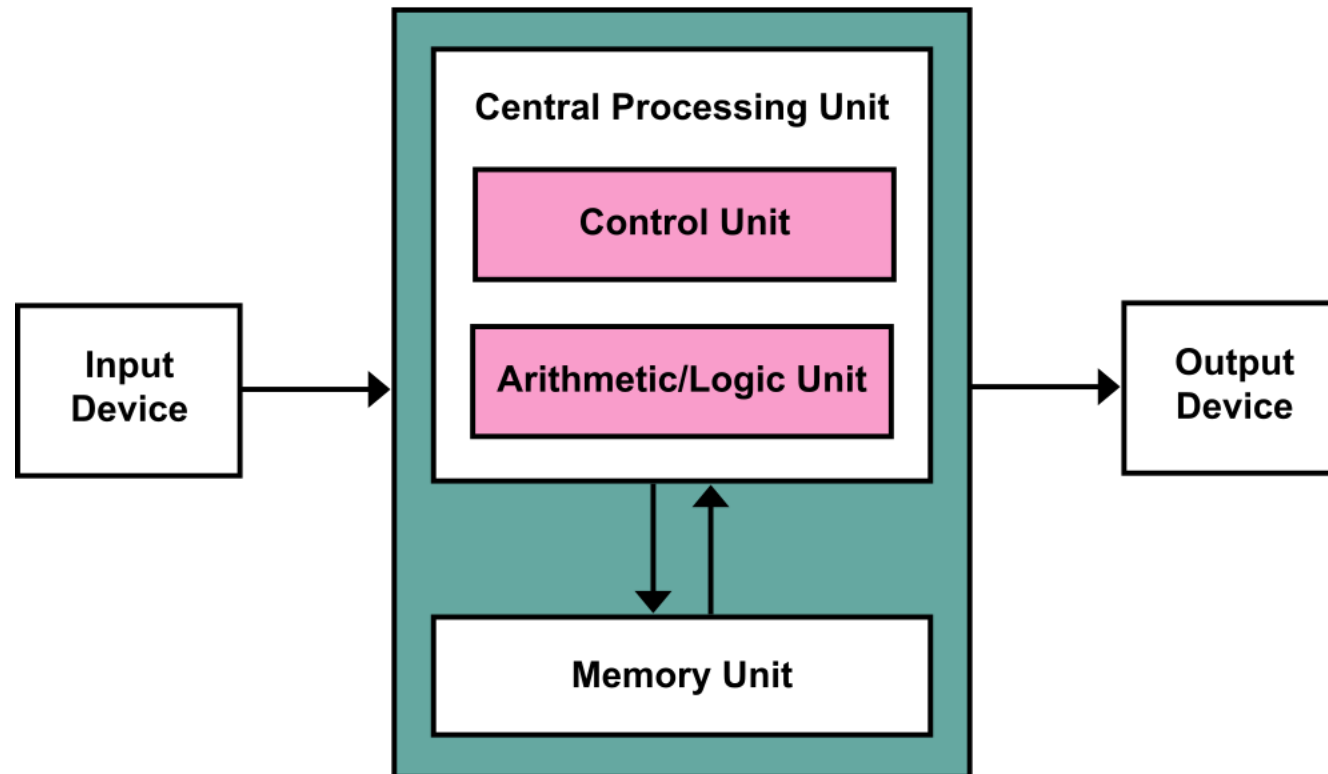
```
void easy_halting2(int x) {  
    while(1);  
}
```

```
void hard_halting(int x) {  
    while (x != 0) {  
        x = (x * (x + 1)) % 73;  
    }  
}
```

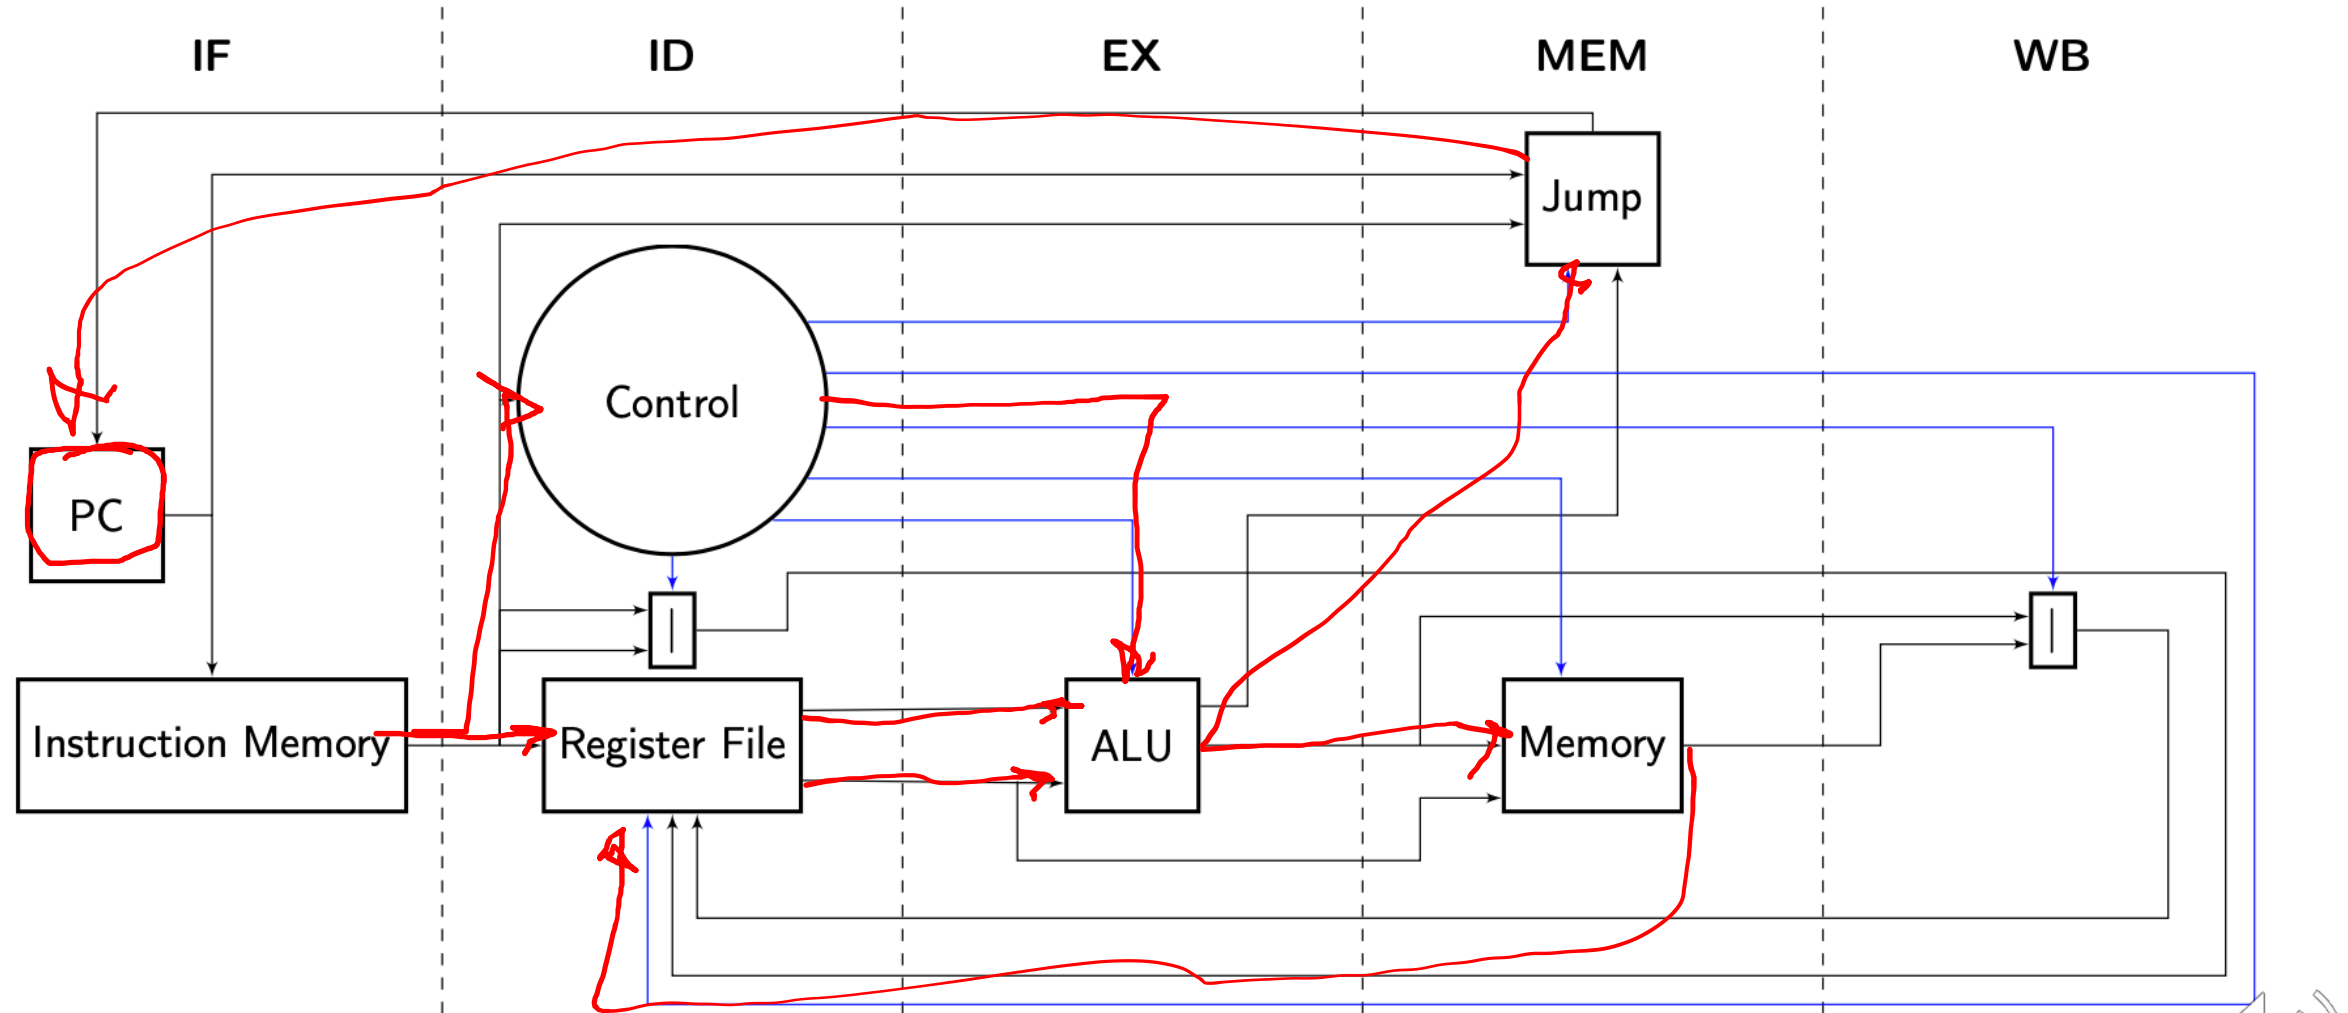
Trying 0 - 0 rounds
Trying 1 - 12 rounds
Trying 2 - 11 rounds
...



The von Neumann Architecture

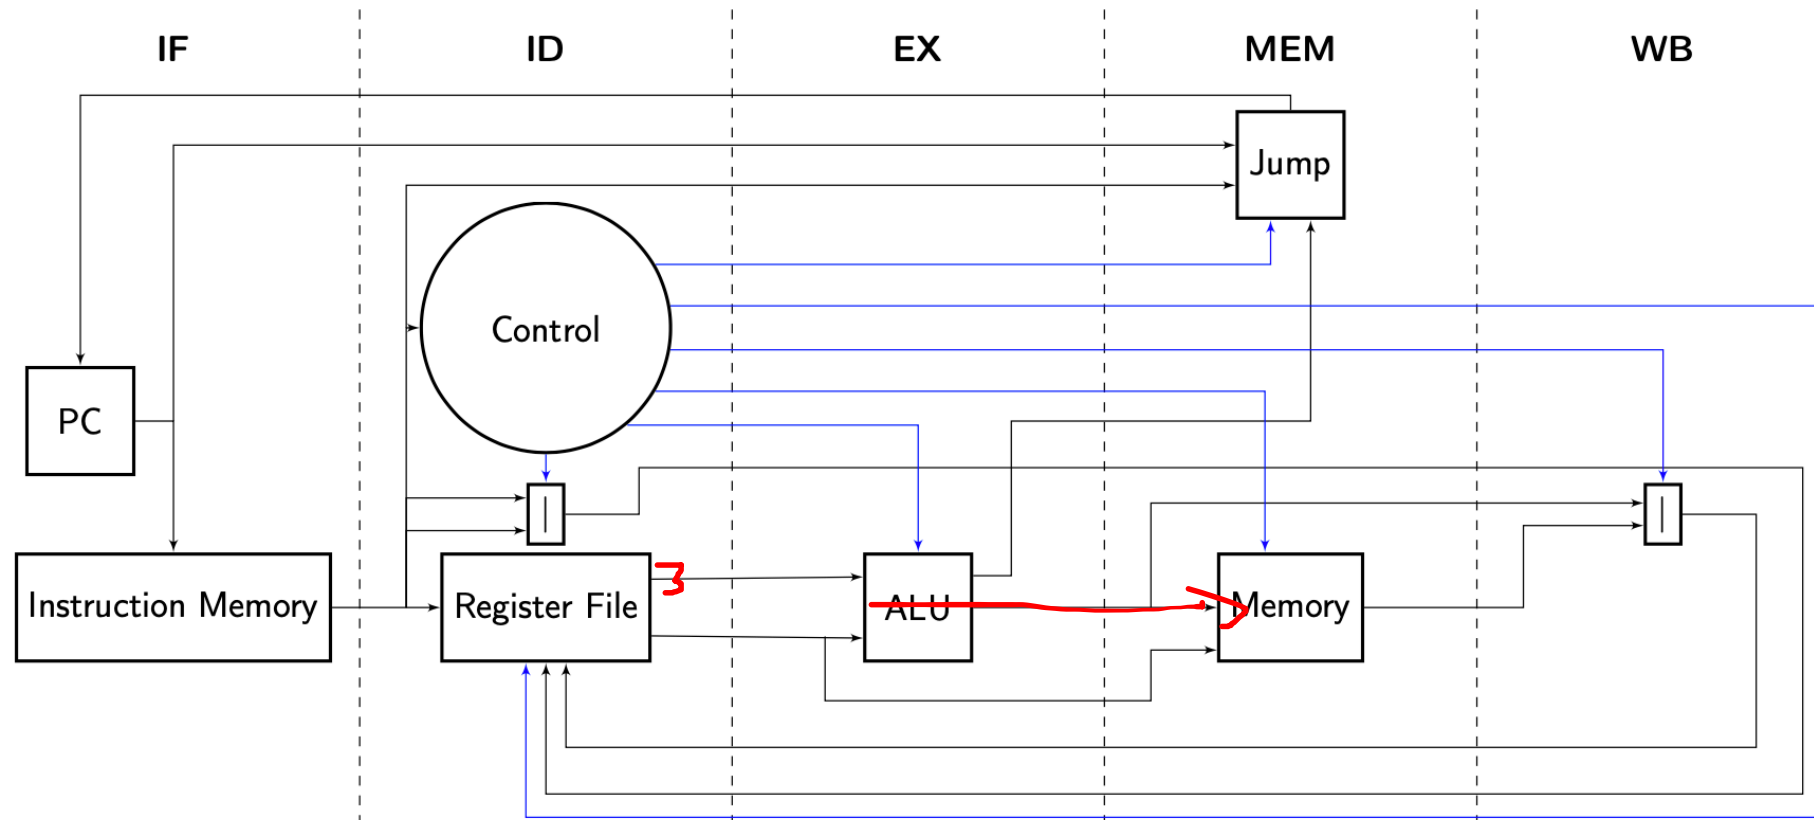


A simple CPU



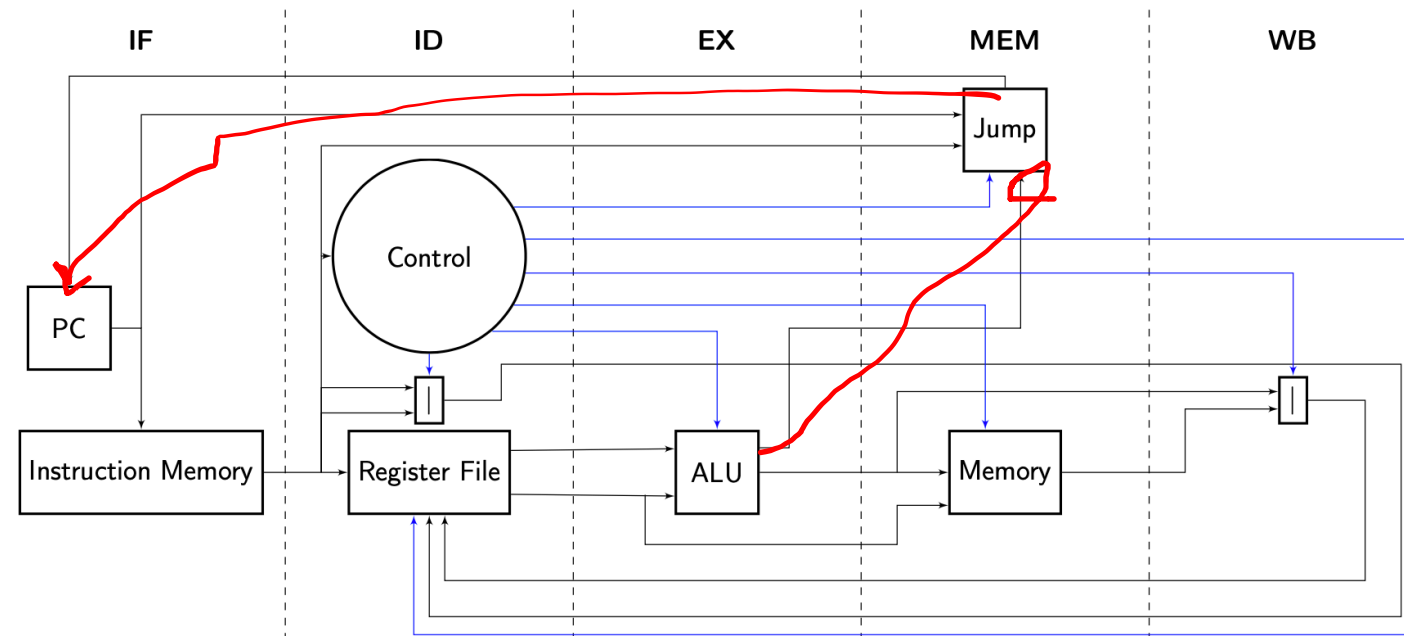
A simple CPU program

```
LOAD_MEM #1234, $1 //Load from memory location 1234 into register 1
LOAD_CONST 1, $2    // Load the value 1 into register 2
ADD $1, $2, $3      // Add register 1 with register 2 and store in register 3
WRITE_MEM $3, #1234 // Write value in register 3 into memory location 1234
```



A simple CPU program with a loop

```
LOAD_MEM #1234, $1    // Load from memory location 1234 into register 1
COPY $1, $2           // Copy the value from register 1 into register 2
LOAD_C 4, $3          // Load the value 3 into register 2
ADD $1, $2, $2        // Add register 1 with register 2 and store in register 2
LOAD_C 1, $4          // Load the value 1 into register $4
SUB $3, $4, $3        // Subtract register 4 from register 3 and store in register 3
JNZ $3, -3            // Jump three instructions back, if register 3 is not zero
WRITE_MEM $2, #1234   // Write value in register 2 into memory location 1234
```



Try it yourself – Modify the program

```
LOAD_MEM #1234, $1    // Load from memory location 1234 into register 1
COPY $1, $2           // Copy the value from register 1 into register 2
LOAD_C 4, $3          // Load the value 3 into register 3
ADD $1, $2, $2        // Add register 1 with register 2 and store in register 2
LOAD_C 1, $4          // Load the value 1 into register $4
SUB $3, $4, $3        // Subtract register 4 from register 3 and store in register 3
JNZ $3, -3            // Jump three instructions back, if register 3 is not zero
WRITE_MEM $2, #1234   // Write value in register 2 into memory location 1234
```

Modify the above program to allow multiplication with a value read from memory



Try it yourself – Modify the program

```
LOAD_MEM #1234, $1    // Load from memory location 1234 into register 1
COPY $1, $2           // Copy the value from register 1 into register 2
LOAD_MEM #1235, $3    // Load from memory location 1234 into register 3
ADD $1, $2, $2        // Add register 1 with register 2 and store in register 2
LOAD_C 1, $4          // Load the value 1 into register $4
SUB $3, $4, $3        // Subtract register 4 from register 3 and store in register 3
JNZ $3, -3            // Jump three instructions back, if register 3 is not zero
WRITE_MEM $2, #1234   // Write value in register 2 into memory location 1234
```

Modify the above program to allow multiplication with a value read from memory



The speed of a CPU

Intel® Core™ i7-3960X Processor Die Detail




Image from Splave


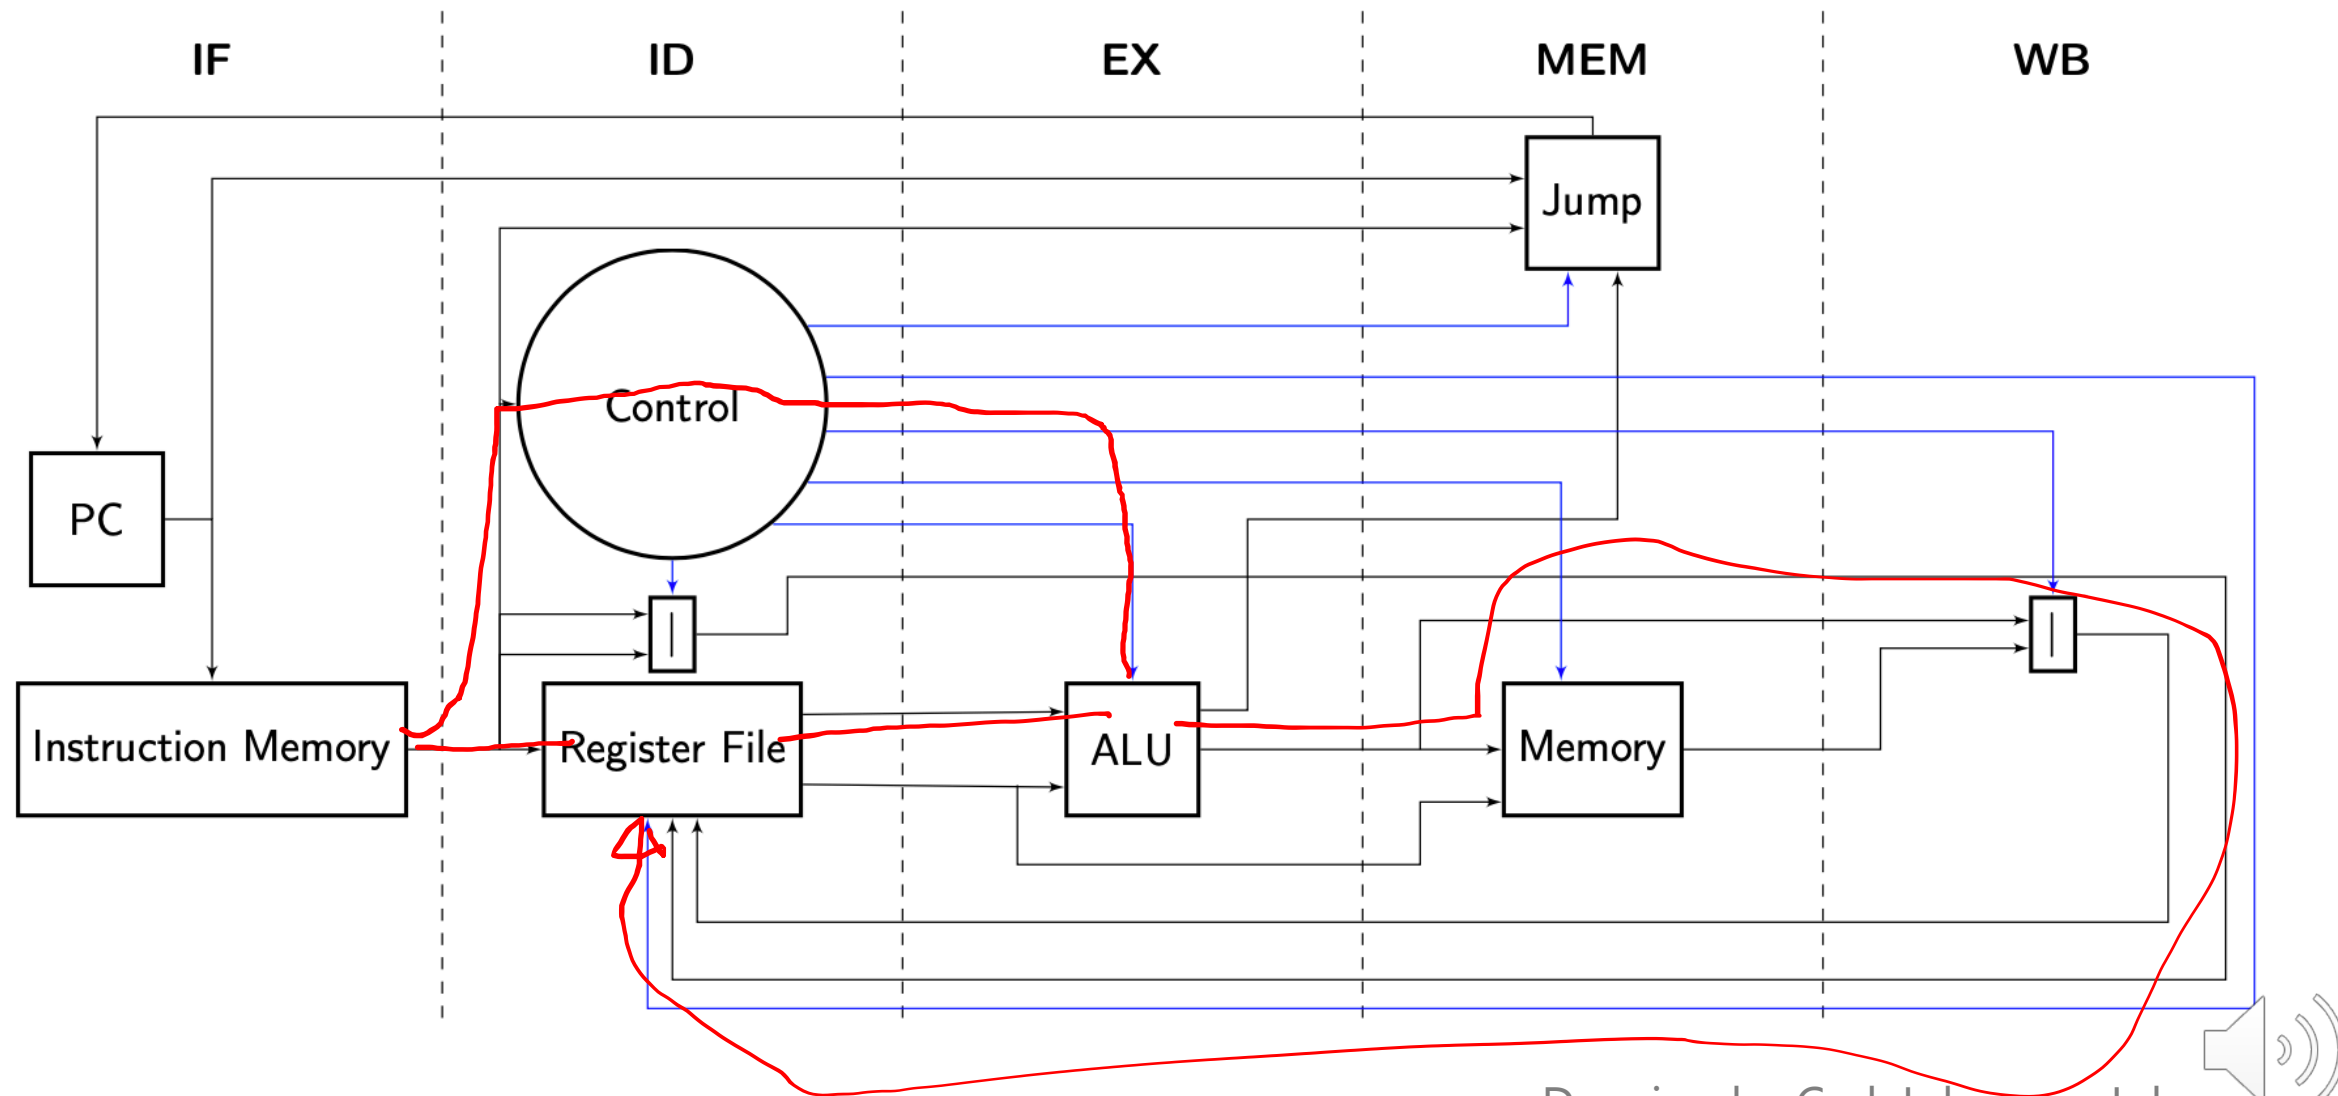


Image from Intel

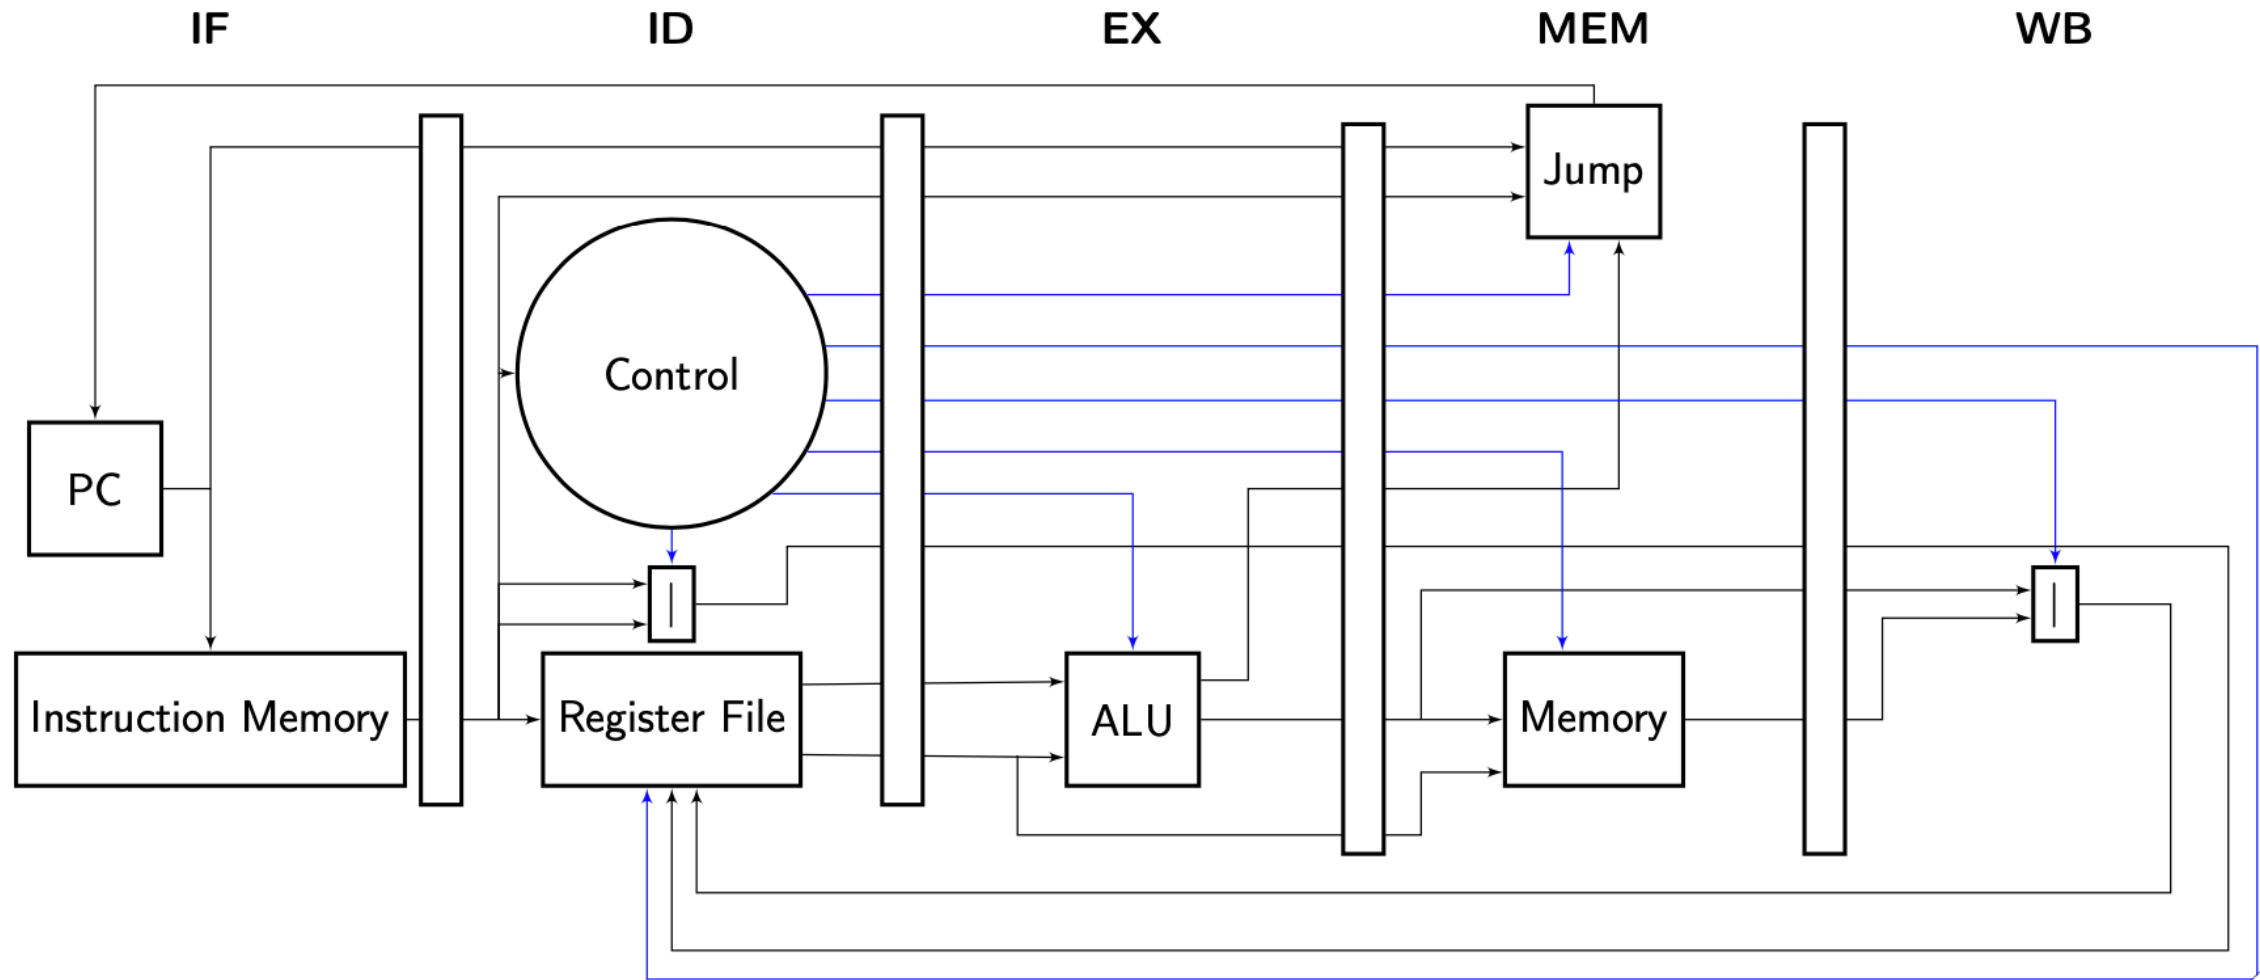


A simple CPU





Drawing by Carl-Johannes Johnsen

A pipelined CPU



Pipelining – Useful in practice



Images from <https://project-navi.blogspot.com/2013/04/pipelining.html>



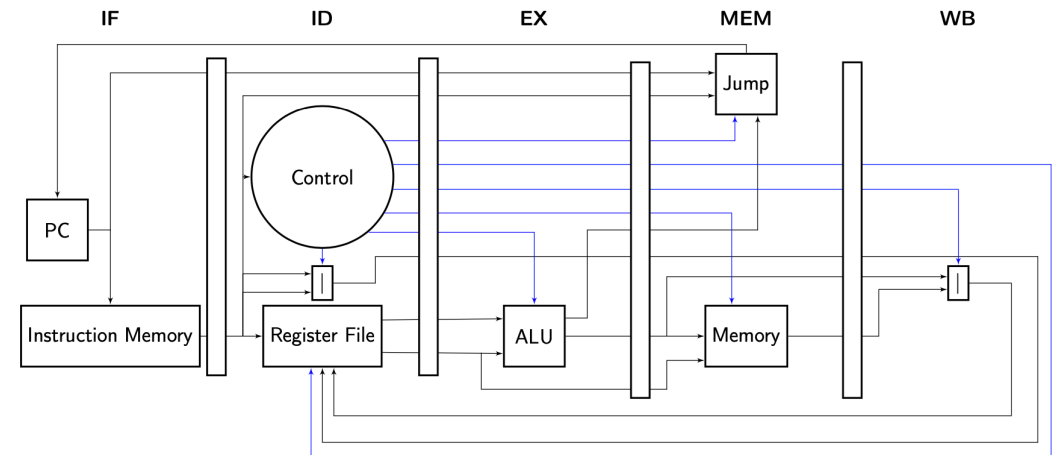
A pipelined CPU – Data hazards

LOAD_MEM #1234, \$1

LOAD_CONST 1, \$2

ADD \$1, \$2, \$3

WRITE_MEM \$3, #1234



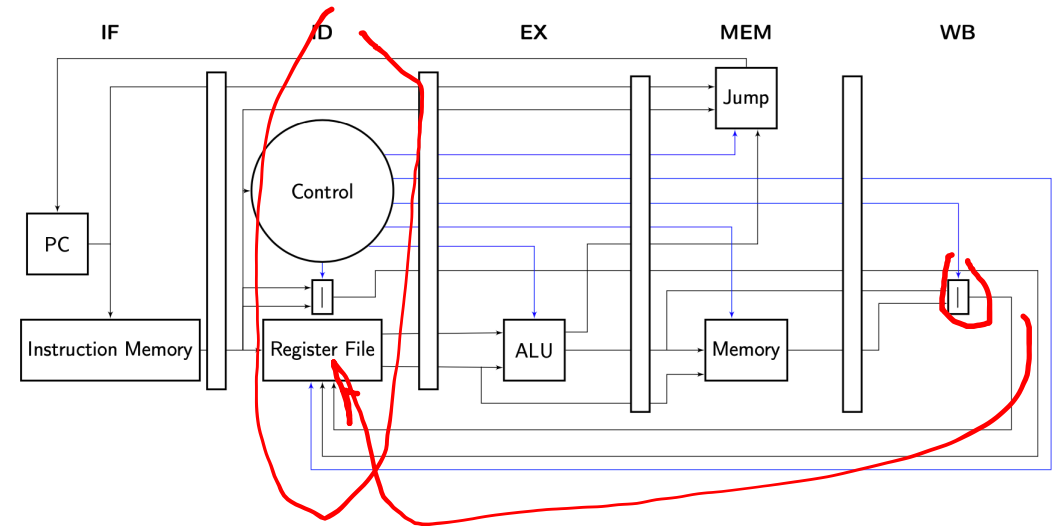
Cycle	IF	ID	EX	MEM	WB
0	LOAD, \$1				
1	LOADC, \$2	LOAD \$1			
2	ADD \$1, \$2, \$3	LOADC, \$2	LOAD \$1		
3	WRITE \$3	ADD \$1, \$2, \$3	LOADC, \$2	LOAD \$1	
4		WRITE \$3	ADD \$1, \$2, \$3	LOADC, \$2	LOAD \$1
5			WRITE \$3	ADD \$1, \$2, \$3	LOADC, \$2
6				WRITE \$3	ADD \$1, \$2, \$3
7					WRITE \$3



A pipelined CPU – Data hazards

```

LOAD_MEM #1234, $1
LOAD_CONST 1, $2
ADD $1, $2, $3
WRITE_MEM $3, #1234
  
```



Cycle	IF	ID	EX	MEM	WB
0	LOAD, \$1				
1	LOADC, \$2	LOAD \$1			
2	ADD \$1, \$2, \$3	LOADC, \$2	LOAD \$1		
3	WRITE \$3	ADD \$1, \$2, \$3	LOADC, \$2	LOAD \$1	
4		WRITE \$3	ADD \$1, \$2, \$3	LOADC, \$2	LOAD \$1
5			WRITE \$3	ADD \$1, \$2, \$3	LOADC, \$2
6				WRITE \$3	ADD \$1, \$2, \$3
7					WRITE \$3

We read \$1 in cycle 3

But \$1 is written in cycle 4, and cannot be read before cycle 5

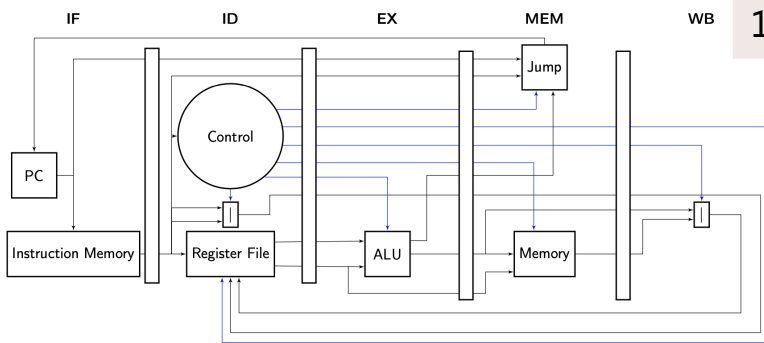


A pipelined CPU – Control hazards

```

LOAD MEM #1234, $1
COPY $1, $2
LOAD C 4, $3
ADD $1, $2, $2
LOAD C 1, $4
SUB $3, $4, $3
JNZ $3, -3
WRITE MEM $2, #1234
  
```

Cycle	IF	ID	EX	MEM	WB
0	LOAD, \$1				
1	COPY \$1, \$2	LOAD \$1			
2	LOADC, \$3	COPY \$1, \$2	LOAD \$1		
3	ADD, \$1, \$2	LOADC, \$3	COPY \$1, \$2	LOAD \$1	
4	LOADC, \$4	ADD, \$1, \$2	LOADC, \$3	COPY \$1, \$2	LOAD \$1
5	SUB, \$3, \$4	LOADC, \$4	ADD, \$1, \$2	LOADC, \$3	COPY \$1, \$2
6	JNZ, \$3	SUB, \$3, \$4	LOADC, \$4	ADD, \$1, \$2	LOADC, \$3
7	WRITE, \$2	JNZ, \$3	SUB, \$3, \$4	LOADC, \$4	ADD, \$1, \$2
8		WRITE, \$2	JNZ, \$3	SUB, \$3, \$4	LOADC, \$4
9			WRITE, \$2	JNZ, \$3	SUB, \$3, \$4
10				WRITE, \$2	JNZ, \$3
11					WRITE, \$2



Do we jump?



RISC or CISC?

CISC – Complex Instruction Set Computer	RISC – Reduced Instruction Set Computer
Hardware focused	Software focused
Multitude of instructions with variations	Limited set of instructions
Fewer registers	More register
Complex addressing modes	Simple addressing modes
Uses microcode (instructions are emulated)	Compiler fixes problems
Instructions can take multiple cycles	Instructions take one cycle
Hard to pipeline	Easy to pipeline



Amdahls Law




Making things faster

$$\text{Speedup} = \frac{\text{execution time before}}{\text{execution time after}}$$

$$\text{Speedup} = \frac{1}{\text{non-enhanced part} + \frac{\text{enhanced part}}{\text{improvement factor}}}$$



Speedup – pipeline example



$$\text{Speedup} = \frac{1}{0.8 + 0.02} = \underline{1.3x}$$



Amdahls Law – Bounds on parallelism


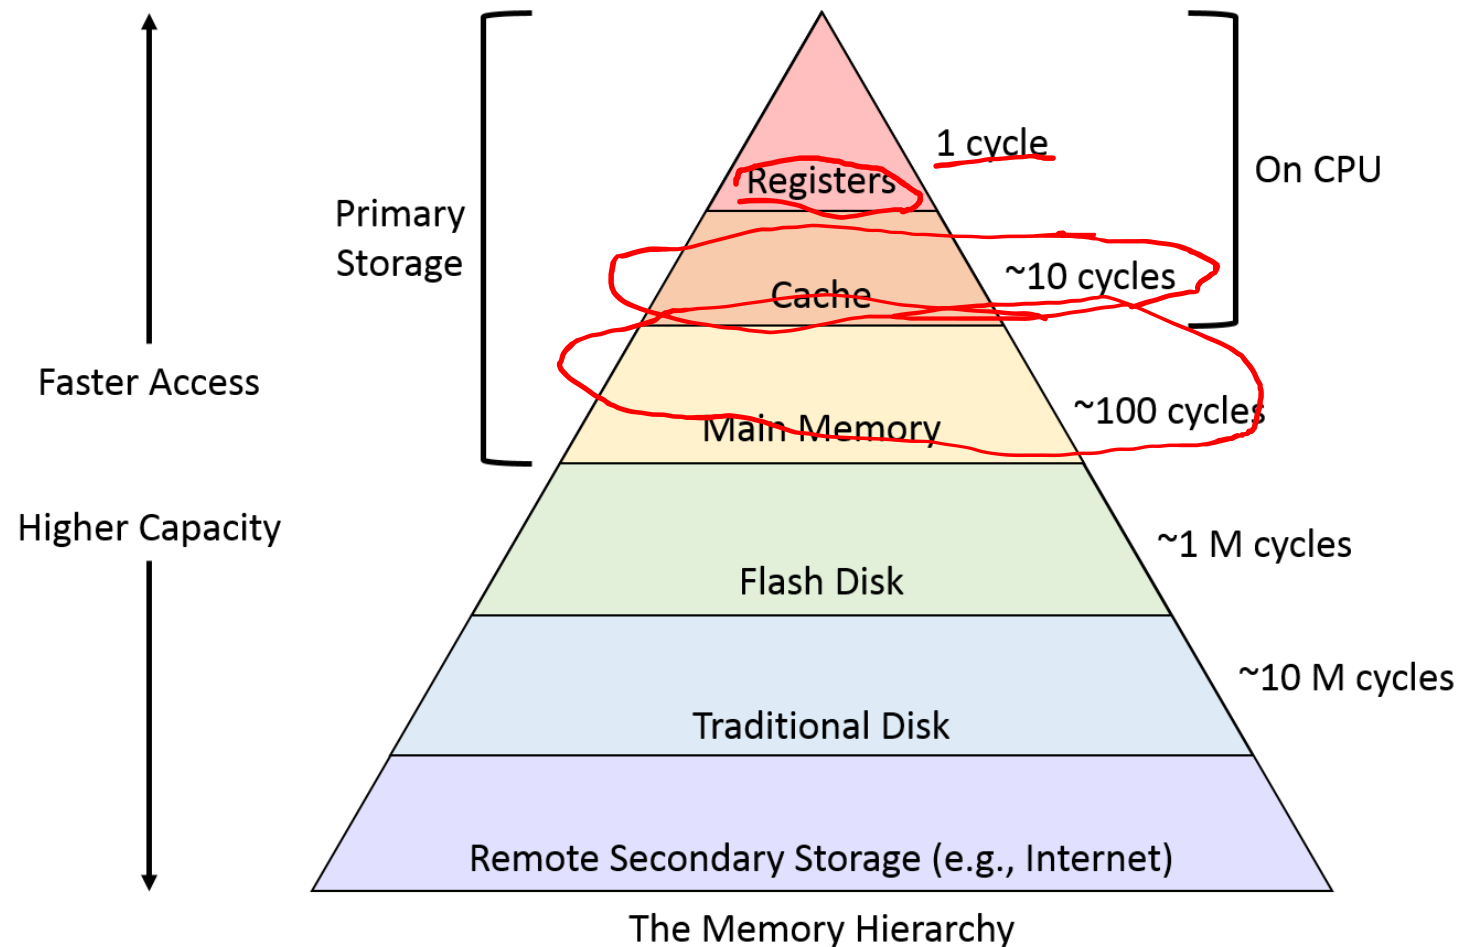


Image from: <https://commons.wikimedia.org/wiki/File:AmdahlsLaw.svg>



Memory – slowing us down



Wrapping up

