

High performance parallel systems

Lecture 2 – Numbers and the C programming language


Kenneth Skovhede, NBI, 2020-11-19

UNIVERSITY OF COPENHAGEN



Recap: Numbers in a digital system

Binary (0b)	Decimal	Octal (0c)	Hexadecimal (0x)
0000	0	0	0
0001	1	1	1
0010	2	2	2
0011	3	3	3
0100	4	4	4
0101	5	5	5
0110	6	6	6
0111	7	7	7
1000	8	10	8
1001	9	11	9
1010	10	12	A
1011	11	13	B
1100	12	14	C
1101	13	15	D
1110	14	16	E
1111	15	17	F
10000	16	20	10



Recap: Addition with binary numbers

$$\begin{array}{r} 111 \\ \underline{0111} (7) \\ + \underline{0011} (3) \\ \hline \underline{1010} (10) \end{array}$$



Recap: Basic gates

AND	0	1
0	0	0
1	0	1

AND



OR




OR	0	1
0	0	1
1	1	1

XOR	0	1
0	0	1
1	1	0



XOR



Not

NOT	0	1
	1	0



Building an adder

$$\underline{A+B = S}$$

+	0	1
0	0	1
1	1	0

Same as XOR

0+0	=	00
0+1	=	01
1+0	=	01
1+1	=	10

Carry	0	1
0	0	0
1	0	1

Same as AND



Building an adder

$$A+B = S$$


+	0	1
0	0	1
1	1	0

Same as XOR

$$\begin{aligned} 0+0 &= 00 \\ 0+1 &= 01 \\ 1+0 &= 01 \\ 1+1 &= 10 \end{aligned}$$

Carry	0	1
0	0	0
1	0	1

Same as AND




$$\begin{array}{r} 111 \\ 0111 \text{ (7)} \\ +0011 \text{ (3)} \\ \hline 1010 \text{ (10)} \end{array}$$



Building an adder

$$A+B = S$$

0+0+0 = 00
0+0+1 = 01
0+1+0 = 01
0+1+1 = 10
1+0+0 = 01
1+0+1 = 10
1+1+0 = 10
1+1+1 = 11




Full-adder

3 inputs
2 outputs



Ripple Carry Adder

$$\begin{array}{r} 111 \\ 0111 \text{ (7)} \\ +0011 \text{ (3)} \\ \hline 1010 \text{ (10)} \end{array}$$



Negative numbers

Bits	Min	Max
4	-8	7
8	-128	127
16	-16384	16383
32	-2147483648	2147483647

Negation is bitwise invert + 1


$-6 \Rightarrow 0b1010$

$\sim 0b1010 + 0b0001 =$

$0b0101 + 0b0001 = 0b0110 = 6$

Odd behavior due to non-symmetry

- $-(-8) = -8$
- $\text{Abs}(-8) = -8$
- $-8 * -1 = -8$



Try it yourself: Addition with Two's complement

Try to compute these basic math problems in binary.
Remember that a negating a number in two's complement
is equivalent to a bitwise inversion + 1

Binary numbers:	Result (in binary)	Decimal
$0b0010 + 0b1010 =$		$2 + -6 = -4$
$0b0101 - 0b1010 =$		
$0b0111 - 0b1000 =$		



Try it yourself: Addition with Two's complement

Try to compute these basic math problems in binary.
Remember that negating a number in two's complement
is equivalent to a bitwise inversion + 1

Binary numbers:	Result (in binary)	Decimal
$0b0010 + 0b1010 =$	$0b1100$	$2 + -6 = -4$
$0b0101 - 0b1010 =$	$0b1011$	$5 - -6 = -5 \text{ (11)}$
$0b0111 - 0b1000 =$	$0b1111$	$7 - -8 = -8 \text{ (15)}$



Multiplication with powers of two

Number	*2 or << 1	*4 or << 2	*8 or << 3
0000 000 <u>1</u> (1)	0000 00 <u>10</u> (2)	0000 0 <u>100</u> (4)	0000 <u>1000</u> (8)
0000 100 <u>1</u> (9)	0001 0010 (18)	0010 0100 (36)	0100 1000 (<u>72</u>)
0001 0110 (22)	0010 1100 (44)	0101 1000 (88)	1011 0000 (176)
1111 000 <u>1</u> (241)	1110 0010 (226)	1100 0100 (196)	1000 1000 (<u>136</u>)
1000 1000 (136)	0001 0000 (16)	0010 0000 (32)	0100 0000 (64)

In binary mode, shifting n bits left is equivalent to multiplying by 2^n

Shifting right is equivalent to dividing by a power of 2

Special instructions are required for negative numbers where the sign bit must be preserved



Real numbers



Fixed point numbers

$$\begin{array}{r} \\ 123.45 \\ +235.56 \\ \hline 359.01 \end{array} \longleftrightarrow \begin{array}{r} \\ 12345 \\ +23556 \\ \hline 35901 \end{array}$$

Currency is often required to use fixed point

Cannot handle ranges, like 1km + 1nm




Floating point numbers

0.0000012345
123450000.0

Despite a large numeric range,
the precision is limited due to the
limited number of bits

One suggestion is to use
two values

Number	Decimal offset
0110 0111	0100



Simpler if we base it
on scientific notation


$$\underline{n * 10^e}$$

But binary: $\underline{n * 2^e}$



Floating point - IEEE-754

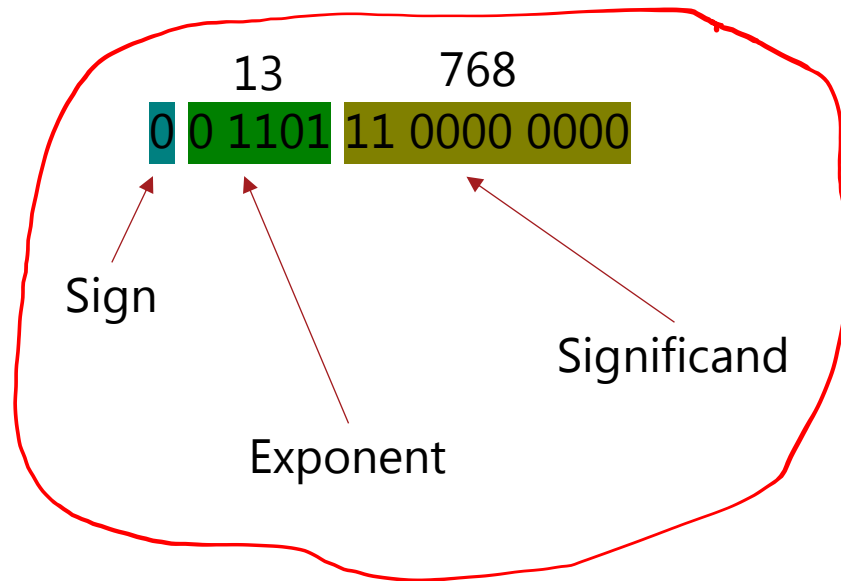
Name	Bits	Exponent bits	Significand bits	Decimal
Half precision	16	5 (-14 to +15)	10(+1) (0 to 1023)	3.31
Single precision	32	8 (-126 to +127)	23(+1)	7.22
Double precision	64	11 (-1022 to +1023)	52(+1)	15.95



For fp32, visit: <https://www.h-schmidt.net/FloatConverter/IEEE754.html>



Floating point - IEEE-754



Actual CPU implementation can differ,
but any stored value must follow this format

$$\text{sign} = 1$$

$$\text{exponent} = 13 - (2^4 - 1) = -2$$


The 11th "bit"

$$\text{value} = \text{sign} * 2^{\text{exponent}} * (1 + \text{significand} / 1024) = 1 * 2^{-2} * (1 + (768/1024)) = 0.4375$$



IEEE – Special numbers

Exponent	Significand	Notes
Not all zero and not all ones	Implicit 1.xxx	Normalized form, most common
All zero	Implicit 0.xxx	Denormalized form, more accuracy for numbers $-1 < n < 1$
All ones	All zero	Infinity, either +inf or -inf
All ones	Non-zero	Not-a-number, NaN



IEEE – Floating point math

$$\begin{array}{r}
 \begin{array}{cc} 13 & 845 \\ \hline 0 & 01101 & 1101001101 \end{array} \\
 + \begin{array}{cc} 13 & 520 \\ \hline 0 & 01101 & 1000001000 \end{array} \\
 \hline
 \end{array}$$

Significand is 2^{10} , meaning [0:1023]

$$= 1 * 2^{(13-15)} * (1 + (845/1024)) = 0.4562988281$$

$$= 1 * 2^{(13-15)} * (1 + (520/1024)) = 0.376953125$$


$$\frac{1365}{1024} = \underline{\underline{0.8332519531}}$$

$$0.8332519531 / 2^{-2} = 3.3330078124$$

$$0.8332519531 / 2^{-1} = \underline{\underline{1.6665039062}}$$

$$\text{Significand} = (1 - 1.6665039062) * 1024 = 683$$


$$= 1 * 2^{-1} * (1 + (683/1024)) = \underline{\underline{0.8334960938}}$$



This means that we have an error of 0.000244140625



IEEE – Floating point - denormalization



$$= 1 * 2^{(8-15)} * (1 + (519/1024)) = 0.01177215576$$


$$= 1 * 2^{(8-15)} * (1 + (520/1024)) = 0.01177978516$$

$$\text{Difference} = -0.000007629394531$$

$$-0.000007629394531 / 2^{-7} = -0.00003051757812$$

$$-0.000007629394531 / 2^{-17} = -1.0$$

Exponent range is [-14, 15]



Denormalized we can describe it as:

$$-1 * 2^{-15} * (0 + 128/1024)$$



Try it yourself: Floating point

$$\begin{array}{c} 20 \qquad 320 \\ 1 \ 1 \ 0100 \ 01 \ 0100 \ 0000 \end{array} = \quad ?$$

$$\begin{array}{c} 16 \qquad 584 \\ 0 \ 1 \ 0000 \ 10 \ 0100 \ 1000 \end{array} = \quad ?$$

$$1/3 = \begin{array}{c} ? \qquad ? \\ ? \ ? \ ???? \ ?? \ ???? \ ???? \end{array}$$

$$12345 = \begin{array}{c} ? \qquad ? \\ ? \ ? \ ???? \ ?? \ ???? \ ???? \end{array}$$

sign = +1 or -1

exponent = $\exp - (2^4 - 1)$

value = $\text{sign} * 2^{\text{exponent}} * (1 + \text{significand} / 1024)$



Try it yourself: Floating point

$$\begin{array}{c} 20 \quad 320 \\ 1 \quad 1 \, 0100 \quad 01 \, 0100 \, 0000 \end{array} = -1 * 2^{(20-15)} * (1 + 320/1024) = -42$$

$$\begin{array}{c} 16 \quad 584 \\ 0 \quad 1 \, 0000 \quad 10 \, 0100 \, 1000 \end{array} = 1 * 2^{(16-15)} * (1 + 584/1024) = \underline{3.140625}$$

Best representation of π with fp16

$$\underline{1/3} = \begin{array}{c} 13 \quad 341 \\ 0 \quad 0 \, 1101 \quad 01 \, 0101 \, 0101 \end{array} = 1 * 2^{(13-15)} * (1 + 341/1024) = \underline{0.3332519531}$$

$$\underline{12345} = \begin{array}{c} 28 \quad 519 \\ 0 \quad 1 \, 1100 \quad 10 \, 0000 \, 0111 \end{array} = 1 * 2^{(28-15)} * (1 + 519/1024) = \underline{12344}$$

sign = +1 or -1

exponent = $\exp - (2^4 - 1)$

value = $\text{sign} * 2^{\text{exponent}} * (1 + \text{significand} / 1024)$



Other ways to compute with real numbers

- Fractions: $1/3 + 1/3 = 2/3$
- Decimal strings "123.7890" + "0.333333333333"
- Symbolic: $\frac{1}{2} * \pi * 4 = 2\pi$



The C programming language



The basics of a C program

```
#include <stdio.h>
```

```
/* Entry function is called main() */
```

```
int main() {  
    // printf() writes to stdout  
    printf("Hello, World!");  
    // Zero means success  
    return 0;
```

```
}
```

#include will copy the file into this file
We need stdio.h to get printf()

/* */ block comments can span multiple lines

Starts the function main() returning an integer

Single line comments end with the line

printf() is the primary output/debug method

return exits the function giving the result back

Blocks in C are written inside curly brackets { ... }

```
#include <stdio.h>
```

```
int main(){printf("Hello, World!");return 0;}
```

Statements end with semicolon ;

Whitespace is ignored (generally)



C – easy to mess up

```
#include <stdio.h>

int main() {
    int b = 2;
    /* some code here */
    if (b < 2)
        //printf("b = %d\n", b);
        b = b * 2;
    return b;
}
```

What is the return value?



C – declaration and definition

```
#include <stdio.h>
```

```
int main() {  
    int a = 1;  
    int b = 2;  
    int c = add(a, b);  
    printf("%d + %d = %d", a, b, c);  
}
```

```
int add(int a, int b) {  
    return a + b;  
}
```

Error: add(int, int) not declared

The method definition



C – declaration and definition

```
#include <stdio.h>

int add(int a, int b);

int main() {
    int a = 1;
    int b = 2;
    int c = add(a, b);
    printf("%d + %d = %d", a, b, c);
}
```

```
int add(int a, int b) {
    return a + b;
}
```

Declare the method

The method is declared, no problems

The method definition



Basic operators in C

Symbol	Description	Example
<code>&</code>	Bitwise AND	<code>1 & 3 = 1</code>
<code> </code>	Bitwise OR	<code>1 2 = 3</code>
<code>^</code>	Bitwise XOR	<code>1 ^ 2 = 3</code>
<code><<</code>	Left shift	<code>1 << 2 = 4</code>
<code>>></code>	Right shift	<code>4 >> 2 = 1</code>
<code>~</code>	Bitwise NOT	<code>~0b0011 = 0b1100</code>

Symbol	Description	Example
<code>&&</code>	Logical AND	<code>True && False = False</code>
<code> </code>	Logical OR	<code>True False = True</code>
<code>!=</code>	Not Equal	<code>True != False = True</code>
<code>!</code>	Logical NOT	<code>!True = False</code>



Beware of logic and bitwise in C

```
int a = 0b0010;  
if (a & 1) // Bitwise  
{  
    // Will not print as 0b0000 is treated as false  
    printf("a & 1 = false\n");  
}  
  
if (a && 1) // Logical  
{  
    // Will print as 0b0010 AND 0b0001 are both treated as true  
    printf("a && 1 = true\n");  
}
```



C – types

Class	Systematic name	Other name	Rank
Integers	Unsigned	_Bool	bool 0
		unsigned char	1
		unsigned short	2
		unsigned int	unsigned 3
		unsigned long	4
		unsigned long long	5
	[Un]signed	char	1
		signed char	1
	Signed	signed short	short 2
		signed int	signed or int 3
		signed long	long 4
		signed long long	long long 5
Floating point		Real	float
	double		
	long double		
	Complex	float _Complex	float complex
		double _Complex	double complex
		long double _Complex	long double complex

Table from "Modern C" by Jens Gustedt



C – types with forced sizes

Signed name	Unsigned name	Bits
int8_t	uint8_t	8
int16_t	uint16_t	16
int32_t	uint32_t	32
int64_t	uint64_t	64



C – platform independence

```
int doubleup(int x) {  
    if (x <= 0)  
        x = 1;  
    else  
        x *= 2;  
    return x;  
}
```

doubleup:
 pushq %rbx # Save caller registers
 subq \$0x18, %rsp # Allocate stack space

 movq \$0, %rbx
 cmp %rax, %rbx # Check input argument
 jg doubleup_nonzero

 movq \$1, %rax # Set to one
 jmp doubleup_done

doubleup_nonzero:
 movq \$2, %rbx # Multiply by 2
 imulq %rbx
doubleup_done:
 addq \$0x18, %rsp # Deallocate stack space
 popq %rbx # Restore registers
 ret # Pop return address and
 # return to caller

x64 assembly



Compiling a program

```
#include <stdio.h>
int main() {
    printf("Hello, World!");
    return 0;
}
```

Contents of file `hello.c`

```
> c99 hello.c
```

Produces a file named `a.out` ... why?

```
> c99 -o hello hello.c
```

Using `-o` lets you pick the filename



Using an IDE to debug

I suggest trying out VS Code (*not* the same as Visual Studio...)

You need to install the C/C++ extension from Microsoft (ms-vscode.cpptools)

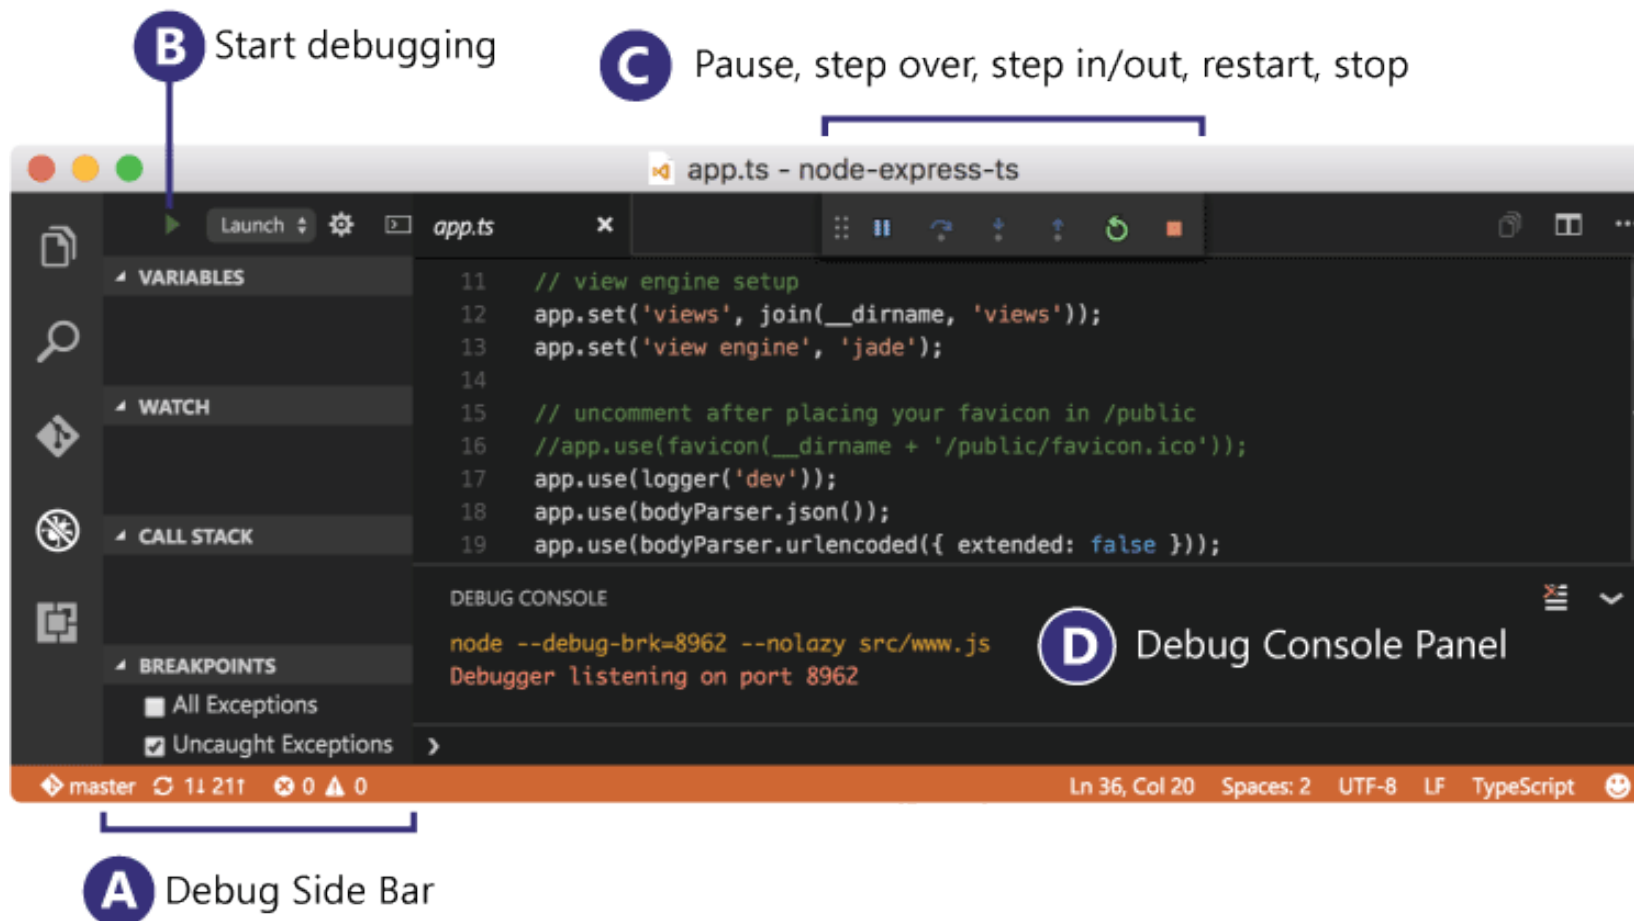


Image from <https://code.visualstudio.com/docs/editor/debugging>




Assignment #1

A floating-point library

aka soft-float

tf116_t



Unbiased, `int5_t`

No special values

No denormalization support



Implement the methods

```
#include <stdlib.h>

typedef uint16_t tfl16_t;

tfl16_t  tfl_sign(tfl16_t value);
int8_t   tfl_exponent(tfl16_t value);
uint16_t tfl_significand(tfl16_t value);
uint8_t  tfl_equals(tfl16_t a, tfl16_t b);
uint8_t  tfl_greaterthan(tfl16_t a, tfl16_t b);
tfl16_t  tfl_normalize(uint8_t sign, int8_t exponent, uint16_t significand);
tfl16_t  tfl_add(tfl16_t a, tfl16_t b);
tfl16_t  tfl_mul(tfl16_t a, tfl16_t b);
```



Wrapping up

