

IO in C

Troels Henriksen

24th of November, 2020

```
printf("Hello , world!\n");
```

```
printf("Hello , world!\n");
```

Hello , world!

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

an integer: 123

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

an integer: 123

```
double y = 1.23;
```

```
printf("a float: %f\n", y);
```



```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;  
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

an integer: 123

```
double y = 1.23;  
printf("a float: %f\n", y);
```

a float: 1.230000

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

an integer: 123

```
double y = 1.23;
```

```
printf("a float: %f\n", y);
```

a float: 1.230000

```
printf("a mess: %d\n", y);
```

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

an integer: 123

```
double y = 1.23;
```

```
printf("a float: %f\n", y);
```

a float: 1.230000

```
printf("a mess: %d\n", y);
```

a mess: 0

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

an integer: 123

```
double y = 1.23;
```

```
printf("a float: %f\n", y);
```

a float: 1.230000

```
printf("a mess: %d\n", y);
```

a mess: 0

Make sure format specifiers and argument types match!

Text representation

- **Machines only understand numbers, and text is an abstraction!**
- E.g. when the terminal receives a byte with the value 65, it draws an A.
- `printf()` determines which *bytes* must be written to the terminal to produce the text corresponding to e.g. the number 123: [49, 50, 51].

Character sets

A character set maps a *number* to a *character*.

- ASCII defines characters in the range 0—127 (asciitable.com).
- Some are invisible/unprintable *control characters*
- *Unicode* is a superset of ASCII that defines tens of thousands of characters for all the world's scripts.

We'll use ASCII, which has the simple property that 1 byte = 1 character.

The ASCII table

Control characters				Normal characters											
000	nul	016	dle	032	□	048	0	064	@	080	P	096	'	112	p
001	soh	017	dc1	033	!	049	1	065	A	081	Q	097	a	113	q
002	stx	018	dc2	034	“	050	2	066	B	082	R	098	b	114	r
003	etx	019	dc3	035	#	051	3	067	C	083	S	099	c	115	s
004	eot	020	dc4	036	\$	052	4	068	D	084	T	100	d	116	t
005	enq	021	nak	037	%	053	5	069	E	085	U	101	e	117	u
006	ack	022	syn	038	&	054	6	070	F	086	V	102	f	118	v
007	bel	023	etb	039	'	055	7	071	G	087	W	103	g	119	w
008	bs	024	can	040	(056	8	072	H	088	X	104	h	120	x
009	tab	025	em	041)	057	9	073	I	089	Y	105	i	121	y
010	lf	026	eof	042	*	058	:	074	J	090	Z	106	j	122	z
011	vt	027	esc	043	+	059	;	075	K	091	[107	k	123	{
012	np	028	fs	044	'	060	<	076	L	092	□	108	l	124	
013	cr	029	gs	045	-	061	=	077	M	093]	109	m	125	}
014	so	030	rs	046	.	062	>	078	N	094	^	110	n	126	~
015	si	031	us	047	/	063	?	079	O	095	_	111	o	127	del

Turning numbers into text

```
int x = 1234;  
printf("x: %d\n", x);
```

Turning numbers into text

```
int x = 1234;  
printf("x: %d\n", x);
```

The text string that is passed to `printf()` looks like this in memory:

Characters	x	:		%	d	\n
Bytes	120	58	32	37	100	10

Turning numbers into text

```
int x = 1234;  
printf("x: %d\n", x);
```

The text string that is passed to `printf()` looks like this in memory:

Characters	x	:		%	d	\n
Bytes	120	58	32	37	100	10

`printf()` rewrites format specifiers (`%d`) to the textual representation of their corresponding value argument:

Characters	x	:		1	2	3	4	\n
Bytes	120	58	32	49	50	51	52	10

These bytes are then written to *standard output* (the terminal) which interprets them as characters and draws pixels on the screen.

Machine representation versus text representation

```
int x = 305419896;
```

- Written as hexadecimal (base-16), this number is 0x12345678.
- One hexadecimal digit is 4 bit, so each group of two digits is one byte, and the number takes four bytes (32 bits).
- The *machine representation* in memory on an x86 CPU is
0x78 0x56 0x34 0x12
- Note that the byte order is *least significant first*—opposite of how we normally write numbers.
- This is because x86 is *little endian*—but *big endian* also exists.
- This *mostly does not matter*, but can be important when we *serialise* machine data to byte sequences for storage in files or transmission over the network.

Writing bytes

The `fwrite` function writes raw data to an open file:

```
size_t fwrite(const void *ptr ,
              size_t size ,
              size_t nmemb,
              FILE *stream );
```

`ptr`: the address in memory of the data.

`size`: the size of each data element in bytes.

`nmemb`: the number of data elements.

`stream`: the target file (opened with `fopen()`).

- Returns the number of data elements written (equal to `nmemb` unless an error occurs).
- Usually no difference between writing one `size x*y` element or `x size-y` elements—do whatever is convenient.

Example of fwrite()

```
#include <stdio.h>

int main() {
    // Open for writing ("w")
    FILE *f = fopen("output", "w");

    char c = 42;

    fwrite(&c, sizeof(char), 1, f);

    fclose(f);
}
```

- Produces a file output.
- File contains the byte 42, corresponding to the ASCII character *.
- **char is just a signed 8-bit integer!**
 - No special “character” meaning.
 - Name is unfortunate/historical.

Another example

```
#include <stdio.h>

int main() {
    FILE *f = fopen("output", "w");

    int x = 0x53505048;
    // Stored as 0x48 0x50 0x50 0x53

    fwrite(&x, sizeof(int), 1, f);

    fclose(f);
}
```

- Writes bytes 0x48 0x50 0x50 0x53
- Corresponds to ASCII characters HPPS
- A big-endian machine would produce SPPH
- **Don't write code that depends on this!**

Converting a non-negative integer to its ASCII representation

```
FILE *f = fopen("output", "w");
int x = 1337;           // Number to write;
char s[10];             // Output buffer.
int i = 10;             // Index of last character written.
while (1) {
    int d = x % 10;      // Pick out last decimal digit.
    x = x / 10;          // Remove last digit.
    i = i - 1;           // Index of next character.
    s[i] = '0' + d;      // Save ASCII character for digit.
    if (x == 0) { break; } // Stop if all digits written.
}
fwrite(&s[i], sizeof(char), 10-i, f); // Write ASCII bytes.
fclose(f);               // Close output file.
```

Reading bytes

```
size_t fread(void *ptr ,  
             size_t size ,  
             size_t nmemb ,  
             FILE *stream );
```

ptr: where to put the data we read.

size: the size of each data element in bytes.

nmemb: the number of data elements.

stream: the target file (opened with `fopen()`).

Very similar to `fwrite()`!

Reading all the bytes in a file

```
#include <stdio.h>
#include <assert.h>

int main(int argc, char* argv[]) {
    FILE *f = fopen(argv[1], "r");
    unsigned char c;
    while (fread(&c, sizeof(char), 1, f) == 1) {
        printf("%3d□", (int)c);
        if (c > 31 && c < 127) {
            fwrite(&c, sizeof(char), 1, stdout);
        }
        printf("\n");
    }
}
```


Running fread-bytes

```
$ gcc -o fread-bytes -Wall -Wextra -pedantic fread-bytes.c
```

Running fread-bytes

```
$ gcc -o fread-bytes -Wall -Wextra -pedantic fread-bytes.c
```

```
$ ./fread-bytes fread-bytes.c
```

```
35 #
```

```
105 i
```

```
110 n
```

```
99 c
```

```
108 l
```

```
117 u
```

```
100 d
```

```
101 e
```

```
32
```

```
60 <
```

```
...
```

Running fread-bytes

```
$ gcc -o fread-bytes -Wall -Wextra -pedantic fread-bytes.c
```

```
$ ./fread-bytes fread-bytes.c
```

```
35 #
```

```
105 i
```

```
110 n
```

```
99 c
```

```
108 l
```

```
117 u
```

```
100 d
```

```
101 e
```

```
32
```

```
60 <
```

```
...
```

```
$ ./fread-bytes fread-bytes
```

```
127
```

```
69 E
```

```
76 L
```

```
70 F
```

```
2
```

```
1
```

```
1
```

```
0
```

```
0
```

```
0
```

```
...
```

Text files versus binary files

- To the system there is no difference between “text files” and “binary files”!
- All files are just byte sequences.
- *Colloquially*: a text file is a file that is understandable when the bytes are interpreted as characters (in ASCII or some other character set).

Text files versus binary files

- To the system there is no difference between “text files” and “binary files”!
- All files are just byte sequences.
- *Colloquially*: a text file is a file that is understandable when the bytes are interpreted as characters (in ASCII or some other character set).

Compactness of storage

- A 32-bit integer takes up to 12 bytes to store as base-10 ASCII digits
- 4 bytes as raw data
- **Raw data takes up less space and is much faster to read.**
- But we need special programs to decode the data to human-readable form.

Conclusions

- Use `printf()` for text output.
- (And `scanf()` for text *input*.)
- Use `fwrite()` to write raw data.
- Use `fread()` to read raw data.
- Raw data files are more compact and faster to read/write.