

# High performance parallel systems


## Lecture 8 – Network protocols and applications

Kenneth Skovhede, NBI, 2020-12-10

UNIVERSITY OF COPENHAGEN



# Client/Server - or Request/Response



Defined in RFC 2616

## HTTP Request

Verb      Path      Protocol version

GET /index.html HTTP/1.1\r\n  
Host: www.example.com\r\n  
User-Agent: my-client-v0.1\r\n  
Cookie: abc=123;def=567\r\n  
X-test: yes!\r\n  
\r\n

Headers

End of header

No body for a GET request

## HTTP Response

Protocol version      Status code and text

HTTP/1.1 200 OK\r\n  
Server: MyWeb/2.2.14\r\n  
Content-Length: 32\r\n  
Content-Type: text/html\r\n  
Connection: Closed\r\n  
\r\n  
<html><body>hello!</body></html>

Body



## Minimal HTTP Request

Verb	Path	Protocol version
GET	/index.html	HTTP/1.1\r\n
Host: www.example.com\r\n		
\r\n		

## Minimal HTTP Response

Protocol version	Status code and text
HTTP/1.1	200 OK\r\n
\r\n	
<u>OK</u>	

More than one request pr. connection:

- Client sends
  - Connection: Keep-Alive
- Server responds
  - Connection: Keep-Alive
  - Keep-Alive: timeout=5, max=50
- All requests and responses must have explicit Content-Length, or boundary marker



# HTTP methods

- GET
  - PUT
  - POST
  - PATCH
  - DELETE
  - (PROPFIND)
- Idempotent operation, caching possible
- Update resource, caching not allowed
- Non-standard operation, used for WebDAV



# The querystring is part of the URL

https://www.google.com/search?q=meaning+of+life

GET /search?q=meaning+of+life HTTP/1.1\r\n  
Host: www.google.com\r\n  
\r\n

Caching allowed, by client, server and proxies

Url format is:

<protocol> :// <server> </path> <?querystring>

Path is referring to the path on the local system but is often restricted to a particular folder.

/

=> /var/www/

/search

=> /var/www/search

/profile/data/set1.txt

=> /var/www/profile/data/set1.txt

/..

=> /var/

Usually forbidden



# Updating a shopping cart - POST

```
POST /cart/add HTTP/1.1
Host: api.example.com
Content-Type: application/x-www-form-urlencoded
Content-Length: 34
```

```
Pizza+with%20cheese=1&salad+bowl=2
```

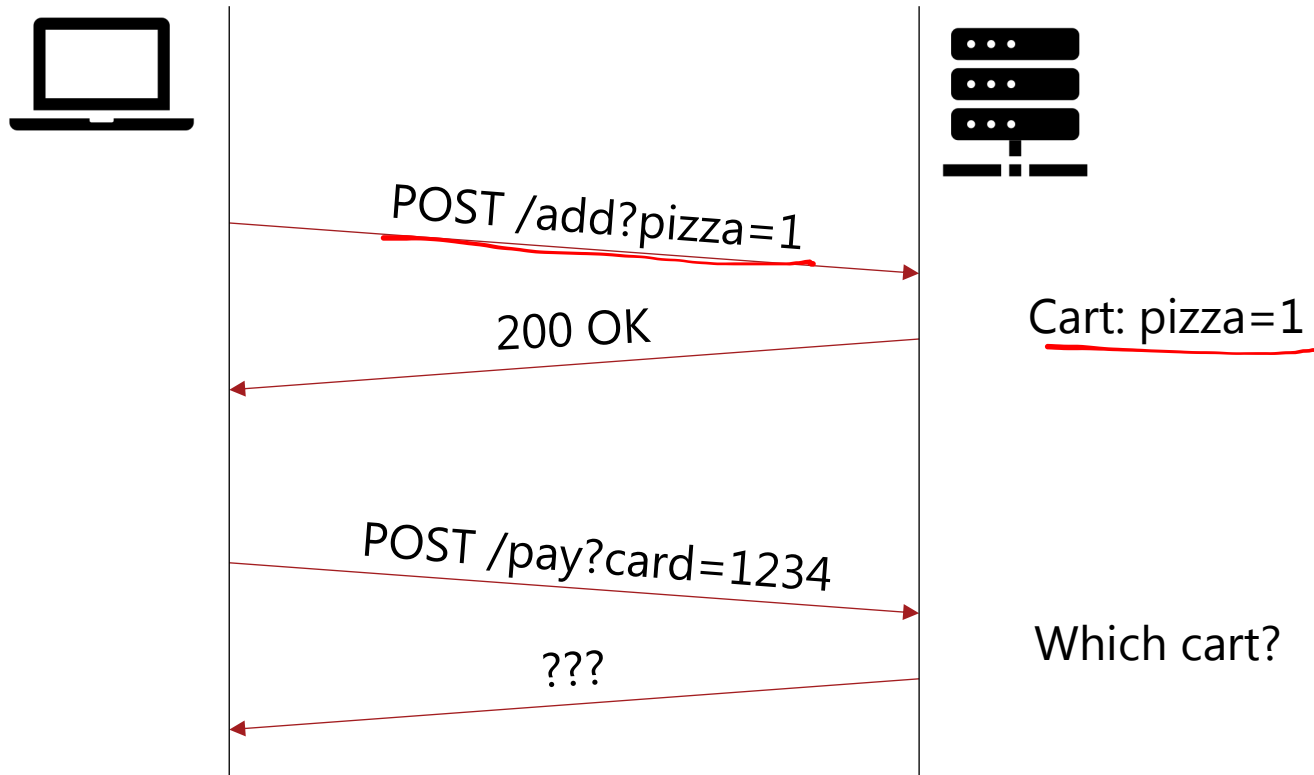
**OR**

```
POST /cart/add HTTP/1.1
Host: api.example.com
Content-Type: multipart/form-data; boundary="boundary"
```

```
--boundary
Content-Disposition: form-data; name="pizza with cheese"
1
--boundary
Content-Disposition: form-data; name="salad bowl"
2
--boundary--
```



# HTTP request/response is stateless





# Cookies for state management

## First visit from client

```
GET /cart HTTP/1.1\r\nHost: api.example.com\r\n\r\n
```

1.

Server creates session

```
HTTP/1.1 200 OK\r\nSet-Cookie: sessionId=1234;Max-Age=500\r\n\r\nOK
```

## Subsequent requests carry the cookie

```
POST /cart/add?pizza=1 HTTP/1.1\r\nHost: api.example.com\r\nCookie: sessionId=1234\r\n\r\n
```

2.

Server matches with session

```
HTTP/1.1 200 OK\r\n\r\n\r\nOK
```

```
PUT /cart/add?pizza=2 HTTP/1.1\r\nHost: api.example.com\r\nCookie: sessionId=1234\r\n\r\n
```

3.

Server can refresh (extend) session

```
HTTP/1.1 200 OK\r\nSet-Cookie: sessionId=1234;Max-Age=500\r\n\r\n\r\nOK
```



# HTTP status codes

- 1xx – connection messages, not about the request
  - 101 – switch protocols, e.g. HTTP/2.0
- 2xx – success messages
  - 200 OK
  - 201 Created
  - 204 No content
- 3xx – redirect messages, not completed
  - 300 – redirect (legacy, vague semantics)
  - 301 – moved permanently
  - 302 – found (or moved temporarily)
  - 304 – not modified
- 4xx – client request error
  - 400 – bad request
  - 401 – not authorized
  - 404 – not found
  - 414 – uri too long
- 5xx – server handling error
  - 500 – internal server error



# Sockets – What is it?

Abstraction for communication, meant to mimic file operations

## File

```
handle = open(path);  
read();  
write();  
close();
```

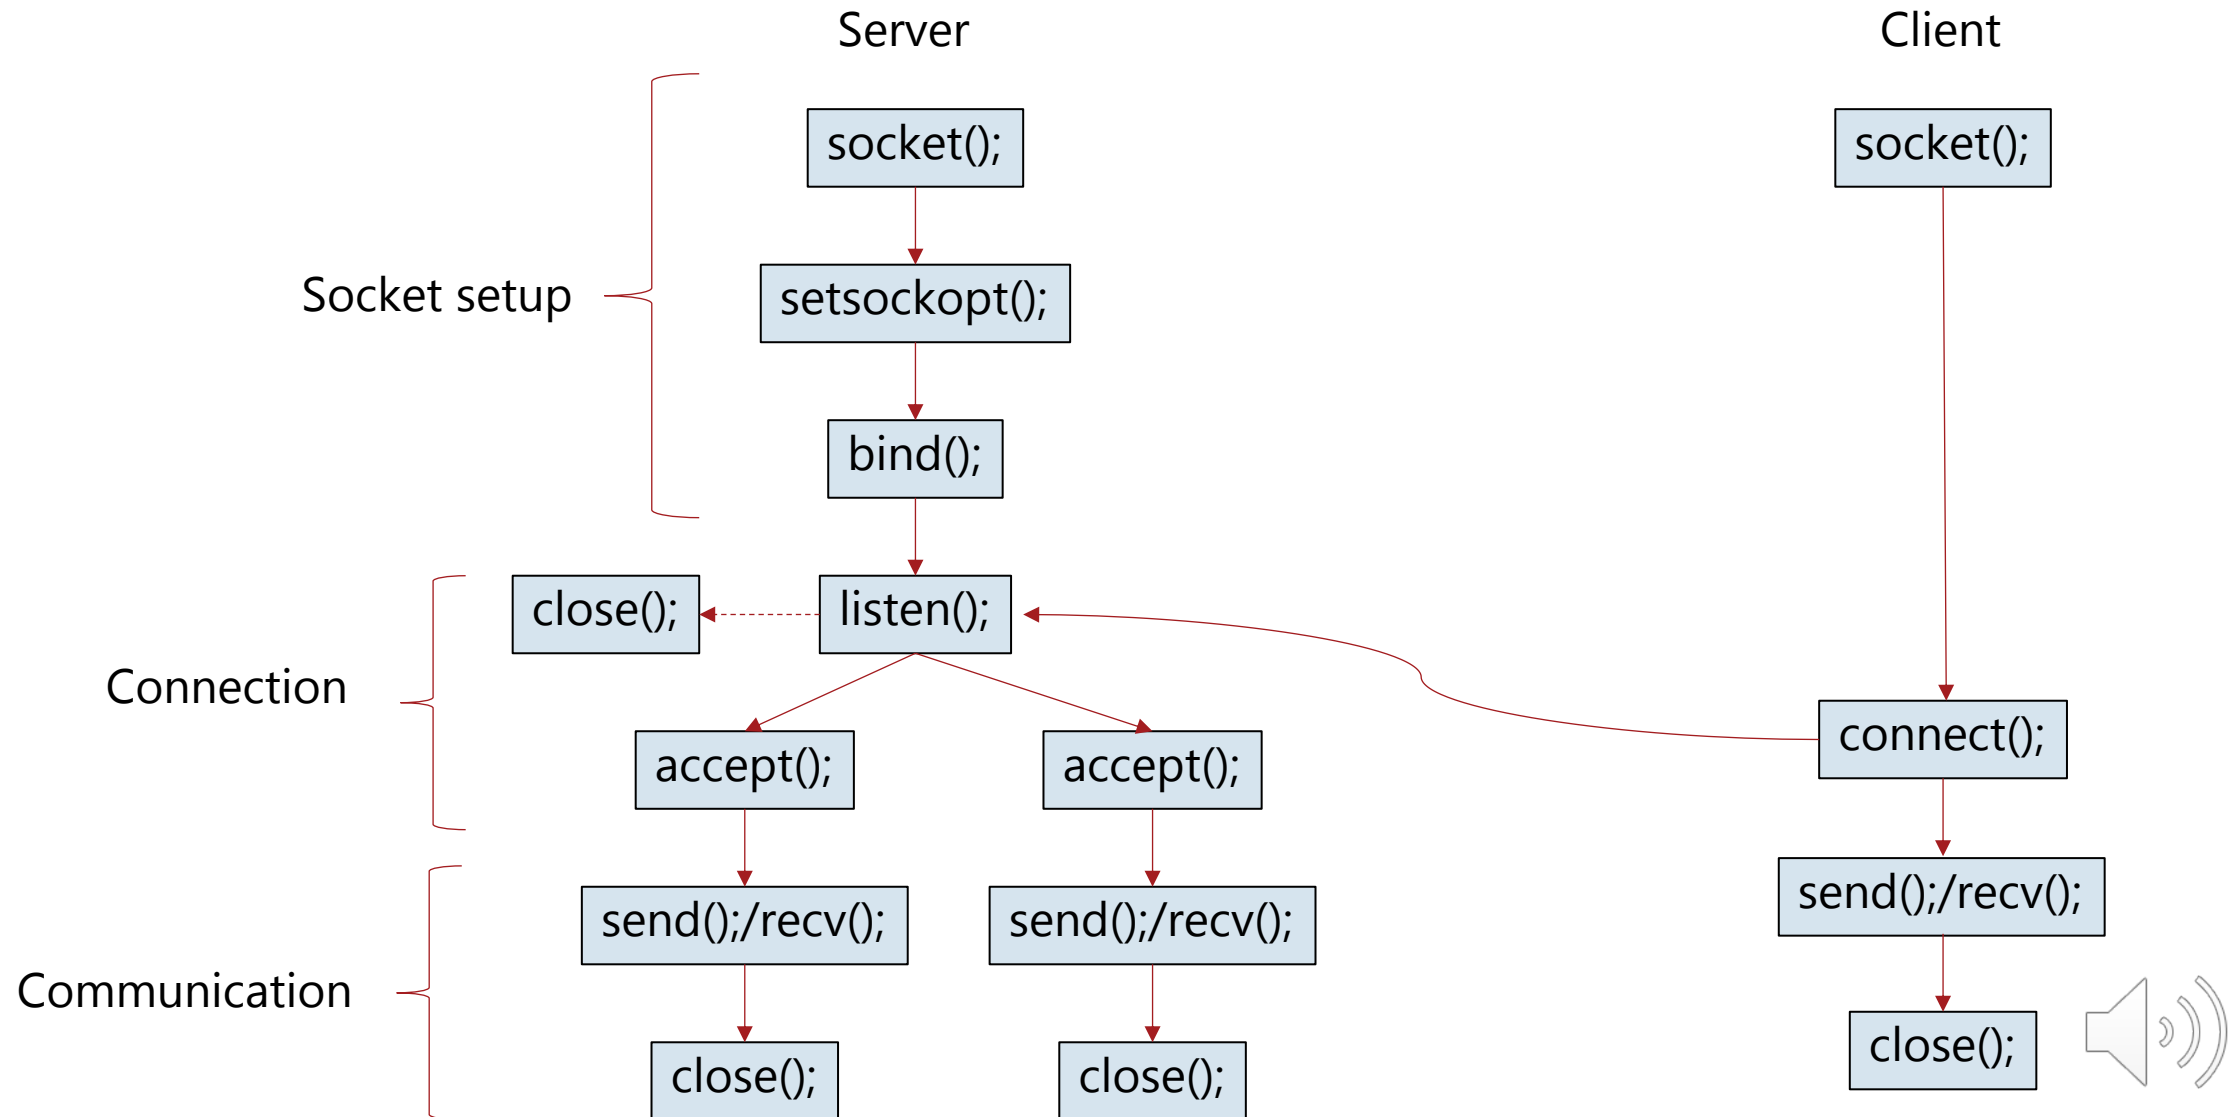
## Socket

```
handle = socket(domain, type, protocol);  
... extra steps here ...  
send();  
recv();  
close();
```

Note that we cannot change position in socket data (i.e. no `seek()`)



# Socket states



# The methods

## In C language

```
int server_fd, new_socket, valread;
struct sockaddr_in address;
char buffer[1024] = {0};
int addrlen = sizeof(address);

// Warning: No error checks!!!
int server_fd = socket(AF_INET, SOCK_STREAM, 0);
address.sin_family = AF_INET;
address.sin_addr.s_addr = INADDR_ANY;
address.sin_port = htons( PORT );
bind(server_fd, (struct sockaddr *)&address, sizeof(address));
listen(server_fd, 3);

new_socket = accept(server_fd,
    (struct sockaddr *)&address, (socklen_t*)&addrlen);
while(1) {
    valread = read( new_socket , buffer, 1024);
    if (valread <= 0)
        break;
    send(new_socket, buffer, valread);
}
```

## In Python

```
import socket
with socket.socket(socket.AF_INET,
    socket.SOCK_STREAM) as s:

    s.bind(('127.0.0.1', 6531))
    s.listen()
    conn, addr = s.accept()
    with conn:
        while True:
            data = conn.recv(1024)
            if not data:
                break
            conn.sendall(data)
```



# ZeroMQ

Extends traditional sockets with more communication models:

- Request/response (like HTTP)
- Publish/subscribe, many-to-many
- Push/pull, producers->consumers->collectors



# ZeroMQ – Publisher in Python

```
import zmq
import random
import sys
import time

port = "5556"
context = zmq.Context()
socket = context.socket(zmq.PUB)
socket.bind("tcp://*:%s" % port)

while True:
    topic = random.randrange(9999,10005)
    messagedata = random.randrange(1,215) - 80
    print "%d %d" % (topic, messagedata)
    socket.send("%d %d" % (topic, messagedata))
    time.sleep(1)
```



# ZeroMQ – Subscriber in Python

```
import sys
import zmq

port = "5556"
context = zmq.Context()
socket = context.socket(zmq.SUB)

print "Collecting updates from server..."
socket.connect ("tcp://localhost:%s" % port)

topicfilter = "10001" # Subscribe to zipcode, default is NYC, 10001
socket.setsockopt(zmq.SUBSCRIBE, topicfilter)

total_value = 0
for update_nbr in range (5):
    string = socket.recv()
    topic, messagedata = string.split()
    total_value += int(messagedata)
    print topic, messagedata

print "Average value for topic '%s' was %dF" % (topicfilter, total_value / update_nbr)
```





# Distributed systems

Has most of the problems from shared memory systems

- But no fast locks

Each exchange must be with messages


- Adds latency to each operation

Machines may crash or go offline at any moment

- Networks can partition forming two or more groups that all consider themselves "global"



# MPI – Message Passing Interface



- No shared memory
- Need to do request/response
- Data can change in between requests
- No total ordering of events



# MPI concepts

- **Communicator** – a “group” for communication
  - Using “**MPI\_COMM\_WORLD**” for every process
- **Size** - the number of processes in the communicator
- **Rank** – The “id” or index of a given process

Point-to-point:

- **mpi\_send()**
- **mpi\_recv()**


Synchronization:

- **mpi\_barrier()**




# MPI – communication


**MPI\_Bcast()**




**MPI\_Scatter()**



**MPI\_Gather()**




**MPI\_Allgather()**




Images from: <https://mpitutorial.com>

# MPI – reduction

**MPI\_Reduce()**



**MPI\_Allreduce()**



Images from: <https://mpitutorial.com>



# MPI with Python

```
from mpi4py import MPI
import numpy

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

if rank == 0:
    data = {'a': 7, 'b': 3.14}
    comm.send(data, dest=1)
elif rank == 1:
    data = comm.recv(source=0)
    print('On process 1, data is ', data)
```



# Peer-to-peer

- Each node is both a client and server
- No pre-defined coordinators nodes
- Handles joining and leaving
- Resources scale with the amount of participants

## Examples

- Skype
- Bitcoin
- Kademlia





# Peer-to-peer example: Kademlia

Basic operation:

- A distributed hash table
- Store key/value pairs
- Retrieve value for given key

Four messages

- PING – liveliness check
- STORE – write a value into the DHT
- FIND\_NODE – locate node with data for a given key
- FIND\_VALUE – locate data for a given key





# Applications on a network

Need to carefully balance workloads

- Automatic balancing systems are great

Avoid any single-point-of-failure

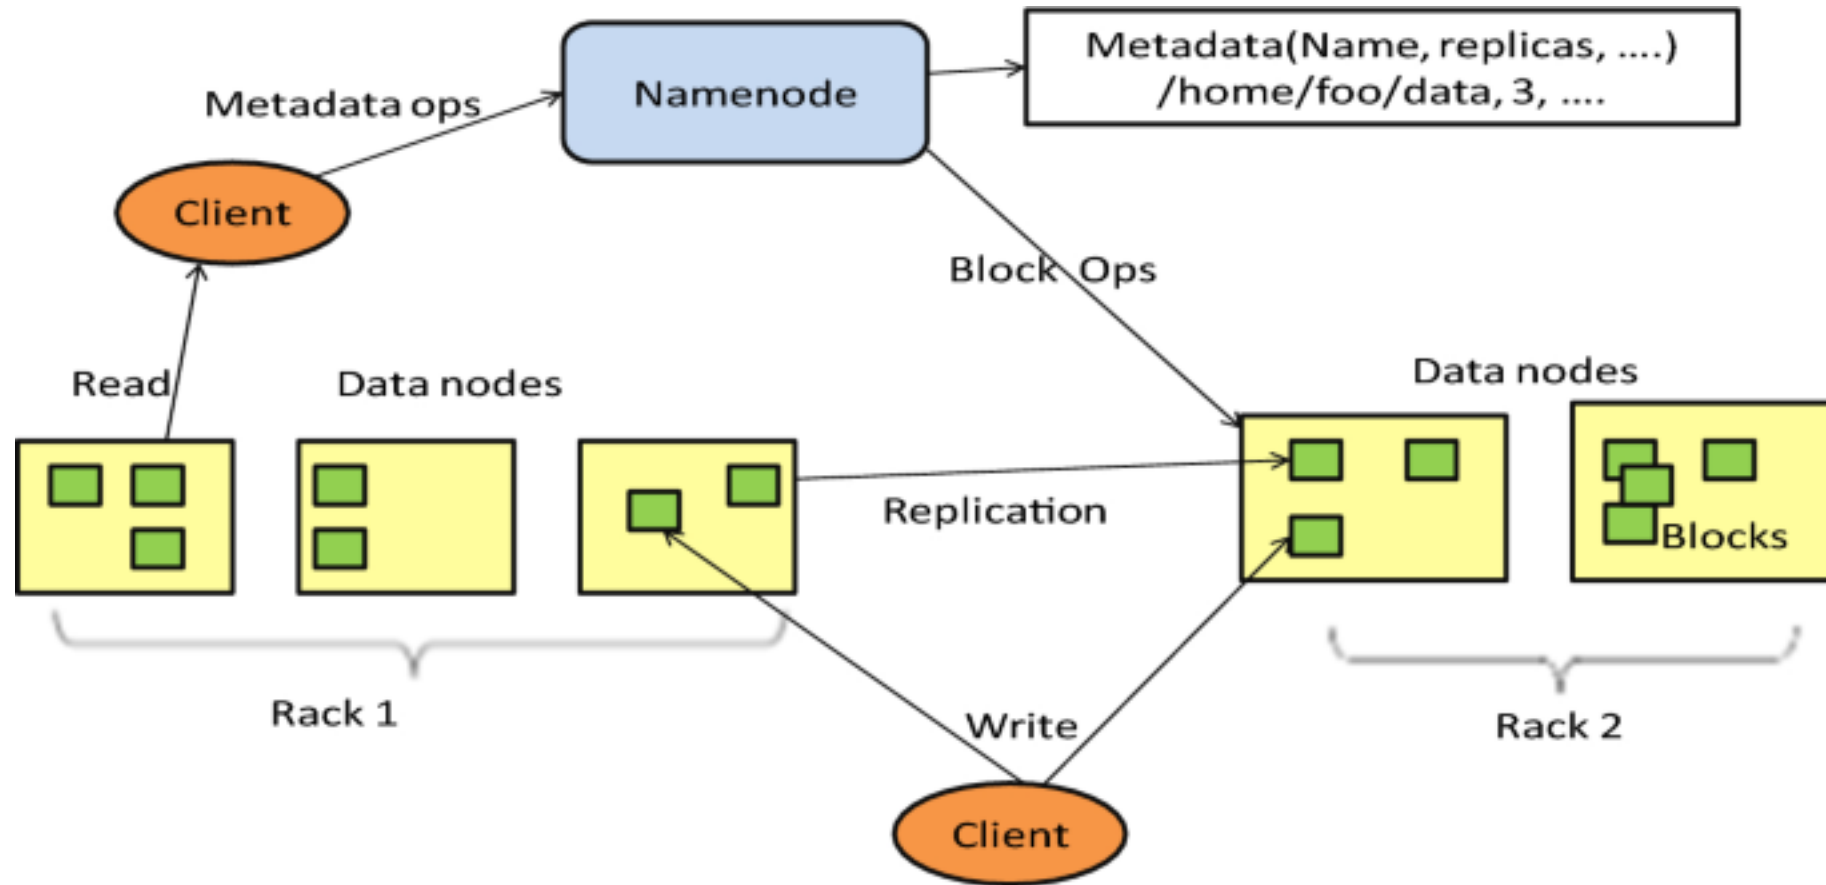
- Really hard in practice

Resilience to duplicate and lost messages

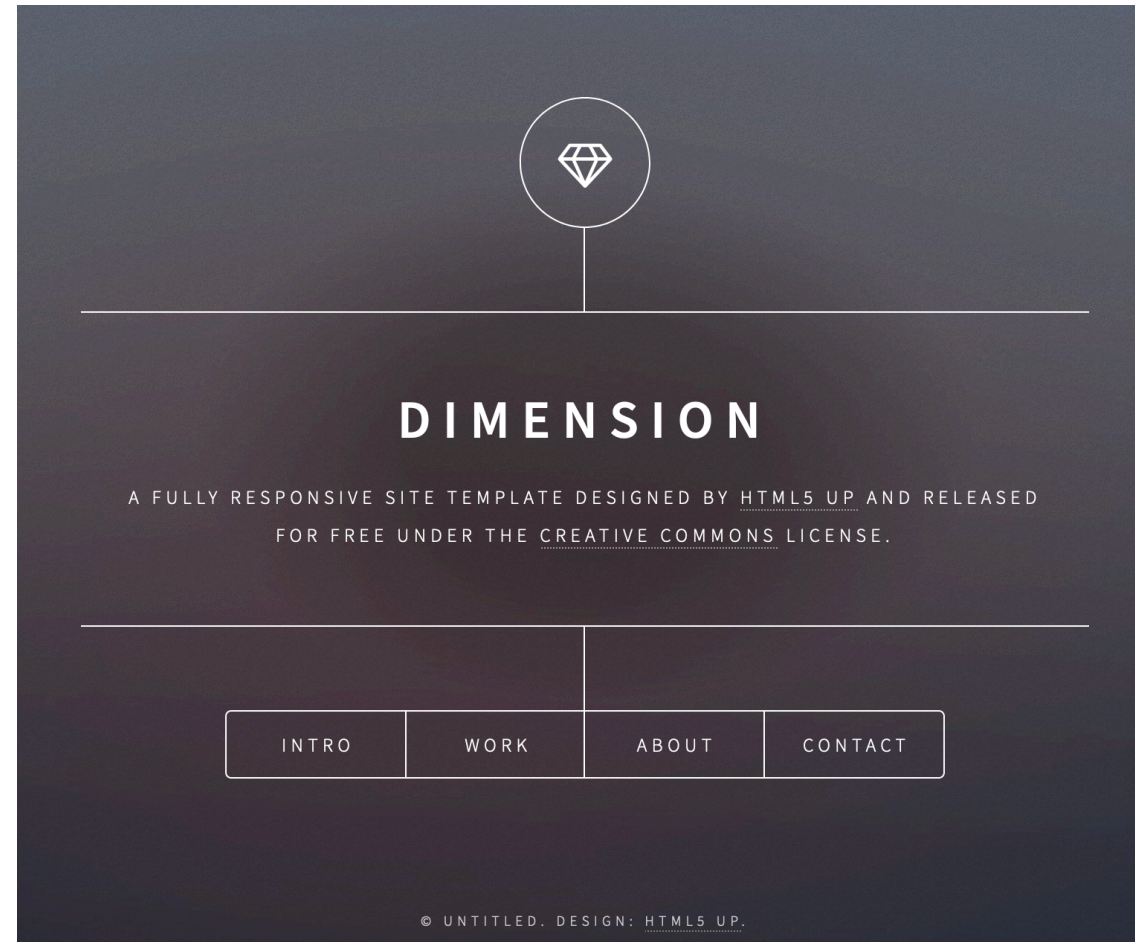
- Almost the same as evil adversary



# Hadoop - HDFS



# Assignment #3 – A web server



Example of a web-page with static files: <https://github.com/cloudacademy/static-website-example>



# Assignment #3 – Suggested tasks

1. Set up a server that listens to requests
  - Can use suggestion from <https://realpython.com/python-sockets/>
2. Write a parser that validates a HTTP/1.1 request
  - Read up on RFC 2616 but *don't read it all!*
  - Be sure to respond correctly to non-GET requests
3. Map URL to local path
  - `GET /images/bg.jpg => /home/user/static-website-example/images/bg.jpg`
  - Beware of tricky paths, eg. `GET ../../passwd`
4. Copy local file to socket
  - May want to set `Content-Type` and `Content-Length` headers

